



Lecture Note 1: OS Introduction

March 1, 2020
Jongmoo Choi

Dept. of software
Dankook University

<http://embedded.dankook.ac.kr/~choijm>

Contents

From Chap 1~2 of the OSTEP

Chap 1. A Dialogue on the Book

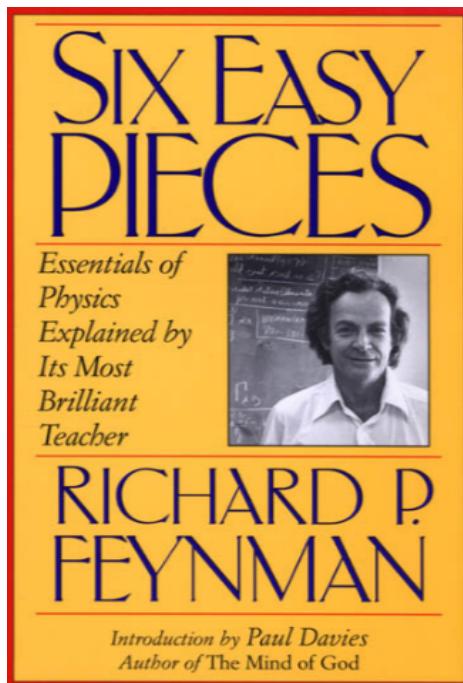
Chap 2. Introduction to Operating System

- ✓ Virtualizing the CPU
- ✓ Virtualizing Memory
- ✓ Concurrency
- ✓ Persistence
- ✓ Design Goals
- ✓ Some History
- ✓ References

Chap 1. A Dialog on the Book

OSTEP

- ✓ Operating Systems: Three Easy Pieces
- ✓ Homage to the Feynman's famous "Six Easy Pieces on Physics"
OS is about half as hard as Physics: from Six to Three Pieces



CONTENTS

vi
Contents

ONE: Atoms in Motion	1	FOUR: Conservation of Energy	69
Introduction	1	What is energy?	69
Matter is made of atoms	4	Gravitational potential energy	72
Atomic processes	10	Kinetic energy	80
Chemical reactions	15	Other forms of energy	81
TWO: Basic Physics	23	FIVE: The Theory of Gravitation	89
Introduction	23	Planetary motions	89
Physics before 1920	27	Kepler's laws	90
Quantum physics	33	Development of dynamics	92
Nuclei and particles	38	Newton's law of gravitation	94
THREE: The Relation of Physics to Other Sciences	47	Universal gravitation	98
Introduction	47	Cavendish's experiment	104
Chemistry	48	What is gravity?	107
Biology	49	Gravity and relativity	112
Astronomy	59	SIX: Quantum Behavior	115
Geology	61	Atomic mechanics	115
Psychology	63	An experiment with bullets	117
How did it get that way?	64	An experiment with waves	120
		An experiment with electrons	122
		The interference of electron waves	124
		Watching the electrons	127
		First principles of quantum mechanics	133
		The uncertainty principle	136
		Index	139

(Source: <https://www.amazon.com/Six-Easy-Pieces-Essentials-Explained/dp/0465025277>)

Chap 1. A Dialog on the Book

OSTEP

~~✗~~

~~✗~~

- ✓ What are Three Pieces: Virtualization, Concurrency, Persistence

Intro	Virtualization		Concurrency	Persistence	Appendices
Preface	3 <u>Dialogue</u>	12 <u>Dialogue</u>	25 <u>Dialogue</u>	35 <u>Dialogue</u>	<u>Dialogue</u>
TOC	4 <u>Processes</u>	13 <u>Address Spaces</u>	26 <u>Concurrency and Threads</u> <small>code</small>	36 <u>I/O Devices</u>	<u>Virtual Machines</u>
1 <u>Dialogue</u>	5 <u>Process API</u> <small>code</small>	14 <u>Memory API</u>	27 <u>Thread API</u>	37 <u>Hard Disk Drives</u>	<u>Dialogue</u>
2 <u>Introduction</u> <small>code</small>	6 <u>Direct Execution</u>	15 <u>Address Translation</u>	28 <u>Locks</u>	38 <u>Redundant Disk Arrays (RAID)</u>	<u>Monitors</u>
	7 <u>CPU Scheduling</u>	16 <u>Segmentation</u>	29 <u>Locked Data Structures</u>	39 <u>Files and Directories</u>	<u>Dialogue</u>
	8 <u>Multi-level Feedback</u>	17 <u>Free Space Management</u>	30 <u>Condition Variables</u>	40 <u>File System Implementation</u>	<u>Lab Tutorial</u>
	9 <u>Lottery Scheduling</u> <small>code</small>	18 <u>Introduction to Paging</u>	31 <u>Semaphores</u>	41 <u>Fast File System (FFS)</u>	<u>Systems Labs</u>
	10 <u>Multi-CPU Scheduling</u>	19 <u>Translation Lookaside Buffers</u>	32 <u>Concurrency Bugs</u>	42 <u>FSCK and Journaling</u>	<u>xv6 Labs</u>
	11 <u>Summary</u>	20 <u>Advanced Page Tables</u>	33 <u>Event-based Concurrency</u>	43 <u>Log-structured File System (LFS)</u>	
		21 <u>Swapping: Mechanisms</u>	34 <u>Summary</u>	44 <u>Flash-based SSDs</u>	
		22 <u>Swapping: Policies</u>		45 <u>Data Integrity and Protection</u>	
		23 <u>Case Study: VAX/VMS</u>		46 <u>Summary</u>	
		24 <u>Summary</u>		47 <u>Dialogue</u>	
				48 <u>Distributed Systems</u>	
				49 <u>Network File System (NFS)</u>	
				50 <u>Andrew File System (AFS)</u>	
				51 <u>Summary</u>	

(Source: <http://pages.cs.wisc.edu/~remzi/OSTEP/>)

Chap 1. A Dialog on the Book

OSTEP

- ✓ What to study?

Professor: *They are the three key ideas we're going to learn about: virtualization, concurrency, and persistence. In learning about these ideas, we'll learn all about how an operating system works, including how it decides what program to run next on a CPU, how it handles memory overload in a virtual memory system, how virtual machine monitors work, how to manage information on disks, and even a little about how to build a distributed system that works when parts have failed. That sort of stuff.*

- ✓ **F** and even a little about how to build a distributed system that works when parts have failed. That sort of stuff.

Student: *I have no idea what you're talking about, really.*

Professor: *Good! That means you are in the right class.*

Student: *I have another question: what's the best way to learn this stuff?*

Professor: *Excellent query! Well, each person needs to figure this out on their own, of course, but here is what I would do: go to class, to hear the professor introduce the material. Then, at the end of every week, read these notes, to help the ideas sink into your head a bit better. Of course, some time later (hint: before the exam!), read the notes again to firm up your knowledge. Of course, your professor will no doubt assign some homeworks and projects, so you should do those; in particular, doing projects where you write real code to solve real problems is the best way to put the ideas within these notes into action. As Confucius said...*

Student: *Oh, I know! 'I hear and I forget. I see and I remember. I do and I understand.' Or something like that.*

But

Chap 2. Introduction to Operating Systems

2.1 Virtualizing CPU

2.2 Virtualizing Memory

2.3 Concurrency

2.4 Persistence

2.5 Design Goals

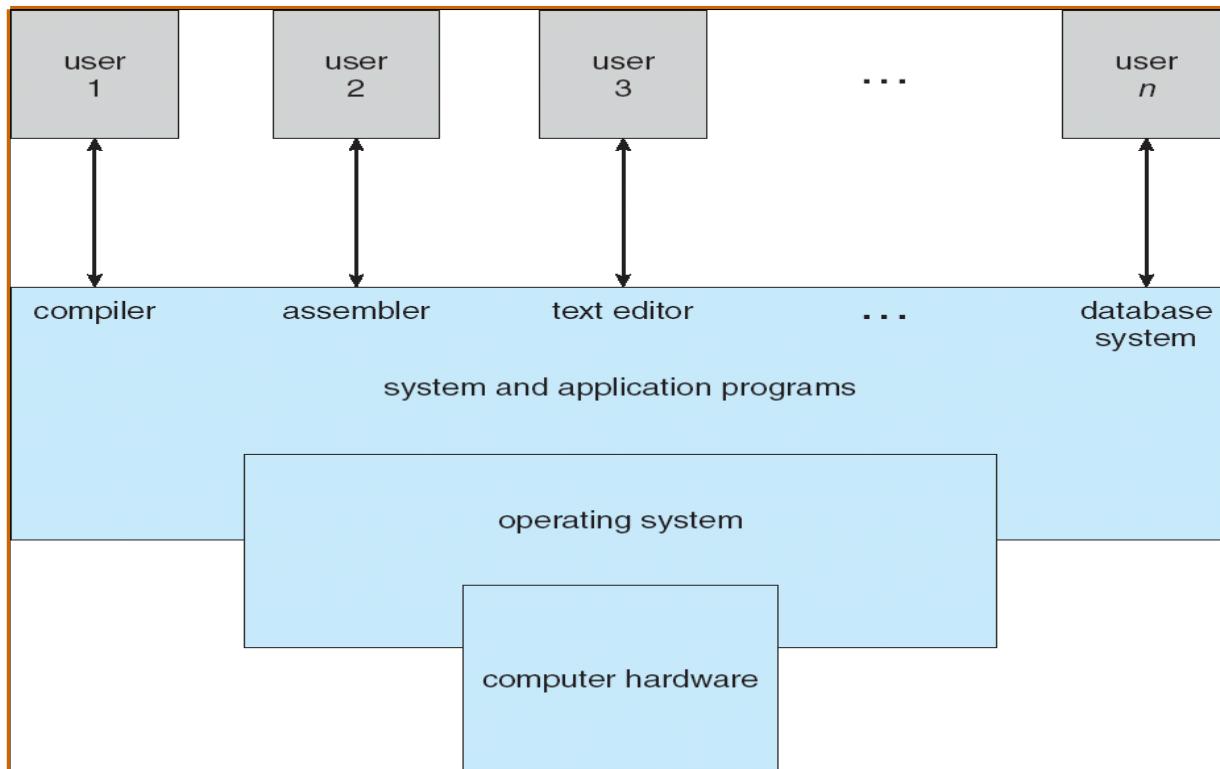
2.6 Some history

2.7 Summary

References

Introduction

Layered structure of a computer system



(Source: A. Silberschatz, “Operating system Concept”)

Introduction

What happens when a program runs?

- ✓ 1. Simple view about running a program

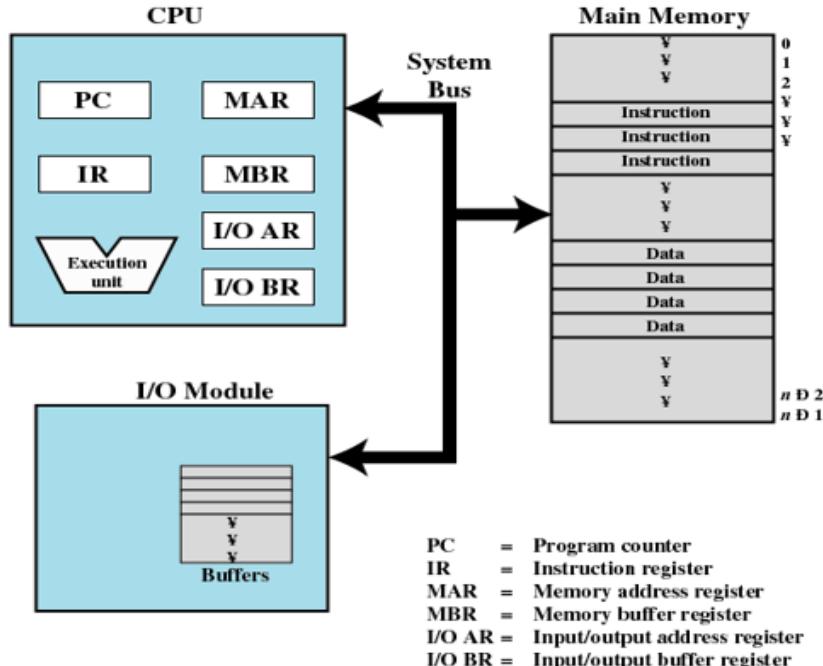


Figure 1.1 Computer Components: Top-Level View

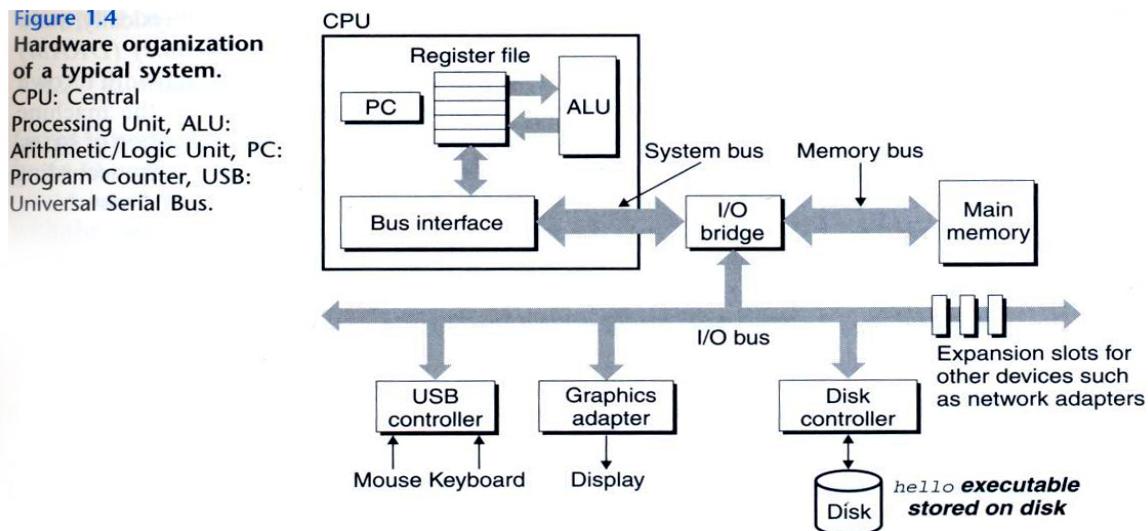
(Source: W. Stallings, "Operating Systems: Internals and Design Principles")

Introduction

What happens when a program runs?

✓ 2. A lot of stuff for running a program

Loading, memory management, scheduling, context switching, I/O processing, file management, IPC, ...
Operating system: 1) make it easy to run programs, 2) operate a system correctly and efficiently



(Source: computer systems: a programmer perspective)

Introduction

Definition of operating system

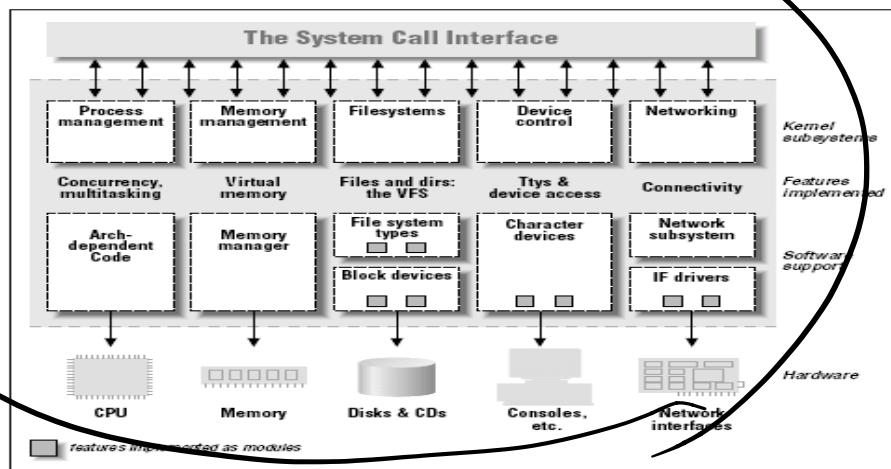
Resource manager

Physical resources: CPU (core), DRAM, Disk, Flash, KBD, Network, ...

Virtual resources: Process, Thread, Virtual memory, Page, File, Directory, Driver, Protocol, Access control, Security, ...

Virtualization (Abstraction)

Transform a physical resource into a more general, powerful, and easy-to-use virtual form



(Source: Linux Device Driver, O'Reilly)

Introduction

System call

- ✓ Interfaces (APIs) provided by OS

	Windows	Unix
Process Control	<code>CreateProcess()</code> <code>ExitProcess()</code> <code>WaitForSingleObject()</code>	<code>fork()</code> <code>exit()</code> <code>wait()</code>
File Manipulation	<code>CreateFile()</code> <code>ReadFile()</code> <code>WriteFile()</code> <code>CloseHandle()</code>	<code>open()</code> <code>read()</code> <code>write()</code> <code>close()</code>
Device Manipulation	<code>SetConsoleMode()</code> <code>ReadConsole()</code> <code>WriteConsole()</code>	<code>ioctl()</code> <code>read()</code> <code>write()</code>
Information Maintenance	<code>GetCurrentProcessID()</code> <code>SetTimer()</code> <code>Sleep()</code>	<code>getpid()</code> <code>alarm()</code> <code>sleep()</code>
Communication	<code>CreatePipe()</code> <code>CreateFileMapping()</code> <code>MapViewOfFile()</code>	<code>pipe()</code> <code>shmget()</code> <code>mmap()</code>
Protection	<code>SetFileSecurity()</code> <code>InitializeSecurityDescriptor()</code> <code>SetSecurityDescriptorGroup()</code>	<code>chmod()</code> <code>umask()</code> <code>chown()</code>

(Source: A. Silberschatz, "Operating system Concept")

Introduction

System call

- ✓ Standard (e.g.. POSIX, Win32, ...)
- ✓ **Mode switch** (user mode, kernel mode)

EXAMPLE OF STANDARD API

As an example of a standard API, consider the `read()` function that is available in UNIX and Linux systems. The API for this function is obtained from the `man` page by invoking the command

```
man read
```

on the command line. A description of this API appears below:

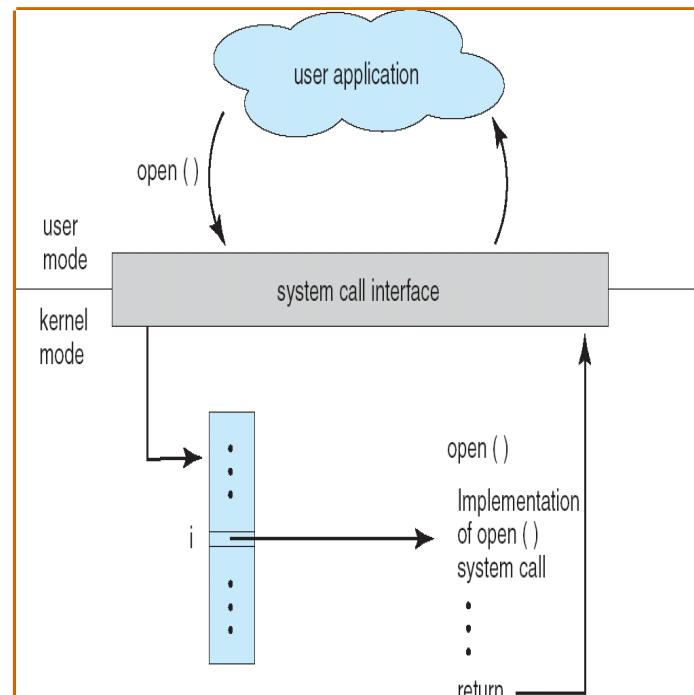
```
#include <unistd.h>
ssize_t read(int fd, void *buf, size_t count)
```

return value	function name	parameters
--------------	---------------	------------

A program that uses the `read()` function must include the `unistd.h` header file, as this file defines the `ssize_t` and `size_t` data types (among other things). The parameters passed to `read()` are as follows:

- `int fd`—the file descriptor to be read
- `void *buf`—a buffer where the data will be read into
- `size_t count`—the maximum number of bytes to be read into the buffer

On a successful read, the number of bytes read is returned. A return value of 0 indicates end of file. If an error occurs, `read()` returns -1.



(Source: A. Silberschatz, "Operating system Concept")

2.1 Virtualizing CPU

A program for the discussion of virtualizing CPU

- ✓ ~~call~~ Spin (busy waiting and return when it has run for a second)
- ✓ print out a string passed in on the command line

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <sys/time.h>
4  #include <assert.h>
5  #include "common.h"
6
7  int
8  main(int argc, char *argv[])
9  {
10     if (argc != 2) {
11         fprintf(stderr, "usage: cpu <string>\n");
12         exit(1);
13     }
14     char *str = argv[1];
15     while (1) { spin(1) 실행될까? 실행될까?
16         Spin(1);
17         printf("%s\n", str);
18     }
19     return 0;
20 }
```

Figure 2.1: Simple Example: Code That Loops and Prints (cpu.c)

2.1 Virtualizing CPU

Execute the CPU program

```
prompt> gcc -o cpu cpu.c -Wall
prompt> ./cpu "A"
A
A
A
A
^C
prompt>
```

Execute

```
prompt> ./cpu A & ./cpu B & ./cpu C & ./cpu D &
[1] 7353
[2] 7354
[3] 7355
[4] 7356
```

A
B
D
C
A
B
D
C
A
...

Process, Scheduling, ...

Figure 2.2: Running Many Programs At Once

2.1 Virtualizing CPU

Issues for Virtualizing CPU

- ✓ How to run a new program?  **process**
- ✓ How to make a new process?  `fork()`
- ✓ How to stop a process?  `exit()`
- ✓ How to execute a new process?  `exec()`
- ✓ How to block a process?  `sleep()`, `pause()`, `lock()`, ...
- ✓ How to select a process to run next?  **scheduling**
- ✓ How to run multiple processes?  **context switch**
- ✓ How to manage multiple cores (CPUs)?  **multi-processor scheduling, cache affinity, load balancing**
- ✓ How to communicate among processes?  **IPC (Inter-Process Communication), socket**
- ✓ How to notify an event to a process?  `signal` (e.g. `^C`)
- ✓ ...



Myth: A process has its own CPU even though there are less CPUs than processes

2.2 Virtualizing Memory

Memory

- ✓ Can be considered as an array of bytes

Another program example

- ✓ Allocate a portion of memory and access it

```
1  #include <unistd.h>
2  #include <stdio.h>
3  #include <stdlib.h>
4  #include "common.h"
5
6  int
7  main(int argc, char *argv[])
8  {
9      int *p = malloc(sizeof(int));
10     assert(p != NULL);
11     printf("(%d) address pointed to by p: %p\n",
12            getpid(), p);
13     *p = 0;
14     while (1) {
15         Spin(1);
16         *p = *p + 1;
17         printf("(%d) p: %d\n", getpid(), *p);
18     }
19     return 0;
20 }
```

Figure 2.3: A Program That Accesses Memory (mem.c)

2.2 Virtualizing Memory

Execute the Mem program

Execute

```
prompt> ./mem
(2134) address pointed to by p: 0x200000
(2134) p: 1
(2134) p: 2
(2134) p: 3
(2134) p: 4
(2134) p: 5
^C
```

proj |

```
prompt> ./mem &; ./mem &
[1] 24113
[2] 24114
(24113) address pointed to by p: 0x200000
(24114) address pointed to by p: 0x200000
(24113) p: 1
(24114) p: 1
(24114) p: 2
(24113) p: 2
(24113) p: 3
(24114) p: 3
(24113) p: 4
(24114) p: 4
...

```

Figure 2.4: Running The Memory Program Multiple Times

Same address but independent

2.2 Virtualizing Memory

Issues for Virtualizing Memory

- ✓ How to manage the address space of a process? ↗ text, data, stack, heap, ...
- ✓ How to allocate memory to a process? ↗ malloc(), calloc(), brk(), ...
- ✓ How to deallocate memory from a process? ↗ free()
- ✓ How to manage free space? ↗ buddy, slab, ...
- ✓ How to protect memory among processes? ↗ virtual memory
- ✓ How to implement virtual memory? ↗ page, segment
- ✓ How to reduce the overhead of virtual memory? ↗ TLB
- ✓ How to share memory among processes? ↗ shared memory
- ✓ How to exploit memory to hide the storage latency? ↗ page cache, buffer cache, ...
- ✓ How to manage NI IMA? ↗ local/remote memory
 - ✗ Any questions? Send an email to me: choijm@dankook.ac.kr
 - ✗ Simple question to take attendance: Explain why the title of our main text include "three easy pieces"? (Hint: see page 3 in this slide, until 6 PM, March 19th)

✗ Illusion: A process has its own unlimited and independent memory even though several processes are sharing limited memory in reality

2.3 Concurrency

Background: how to create a new scheduling entity?

- ✓ Two programming model: process and thread
- ✓ Key difference: data sharing

// fork example (Refer to the Chapter 5 in OSTEP)

// by J. Choi (choijm@dku.edu)

```
#include <stdio.h>
#include <stdlib.h>
```

int a = 10;

void *func()

{

a++;

printf("pid = %d\n", getpid());

}

int main()

{

int pid;

if ((pid = fork()) == 0) { //need exception handle

func();

exit(0);

}

wait();

printf("a = %d by pid = %d\n", a, getpid());



// thread example (Refer to the Chapter 27 in OSTEP)

// by J. Choi (choijm@dku.edu)

```
#include <stdio.h>
#include <stdlib.h>
```

int a = 10;

void *func()

{

a++;

printf("pid = %d\n", getpid());

}

int main()

{

int p_thread;

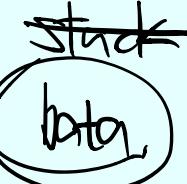
if (pthread_create(&p_thread, NULL, func, (void *)NULL)) < 0

exit(0);

}

pthread_join(p_thread, (void *)NULL);

printf("a = %d by pid = %d\n", a, getpid());



2.3 Concurrency

Concurrency

- ✓ ~~Problems~~ arise when working on many things simultaneously on the same data

A program for discussing concurrency

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include "common.h"
4
5  volatile int counter = 0;
6  int loops;
7
8  void *worker(void *arg) {
9      int i;
10     for (i = 0; i < loops; i++) {
11         counter++;
12     }
13     return NULL;
14 }
15
16 int
17 main(int argc, char *argv[])
18 {
19     if (argc != 2) {
20         fprintf(stderr, "usage: threads <value>\n");
21         exit(1);
22     }
23     loops = atoi(argv[1]);
24     pthread_t p1, p2;
25     printf("Initial value : %d\n", counter);
26
27     Pthread_create(&p1, NULL, worker, NULL);
28     Pthread_create(&p2, NULL, worker, NULL);
29     Pthread_join(p1, NULL);
30     Pthread_join(p2, NULL);
31     printf("Final value   : %d\n", counter);
32
33 }
```

Thread

Figure 2.5: A Multi-threaded Program (threads.c)

2.3 Concurrency

Execute the multi-thread program

```
prompt> gcc -o thread thread.c -Wall -pthread
prompt> ./thread 1000
Initial value : 0
Final value : 2000
```

Programming model

```
✓ prompt> ./thread 100000
Initial value : 0
Final value : 143012 // huh??
✓ prompt> ./thread 100000
Initial value : 0
Final value : 137298 // what the??
```

race loop
control

Why? → Atomicity
→ 531226

⇒ concurrency control이 필요하다.

2.3 Concurrency

Issues for Concurrency

- ✓ How to support concurrency correctly? lock(), semaphore()
- ✓ How to implement atomicity in hardware? test and set(), swap()
- ✓ What is the semaphore?
- ✓ What is the monitor?
- ✓ How to solve the traditional concurrent problems such as **producer-consumer, readers-writers and dining philosophers?**
- ✓ What is a deadlock?
- ✓ How to deal with the deadlock?
- ✓ How to handle the timing bug?
- ✓ What is the asynchronous I/Os?

...

Illusion: Multiple processes run in a cooperative manner on shared resources even though they actually race with each other on the resources

2.4 Persistence

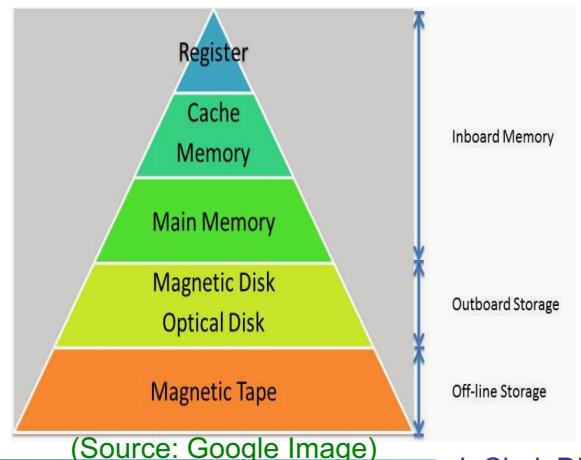
Background: DRAM vs. Disk



VS



- ✓ Capacity, Speed, Cost, ...
- ✓ Access granularity: Byte vs. Sector
- ✓ Durability: Volatile vs. Non-volatile



2.4 Persistence

Persistence

- ✓ Users want to maintain data permanently (durability)
- ✓ DRAM is volatile, requiring write data into storage (disk, SSD) explicitly

A program for discussing persistence

- ✓ Use the notion of a file (not handle disk directly)

```
1 #include <stdio.h>
2 #include <unistd.h>
3 #include <assert.h>
4 #include <fcntl.h>
5 #include <sys/types.h>
6
7 int
8 main(int argc, char *argv[])
9 {
10     int fd = open("/tmp/file", O_WRONLY | O_CREAT | O_TRUNC, S_IRWXU);
11     assert(fd > -1);
12     int rc = write(fd, "hello world\n", 13);
13     assert(rc == 13);
14     close(fd);
15     return 0;
16 }
```

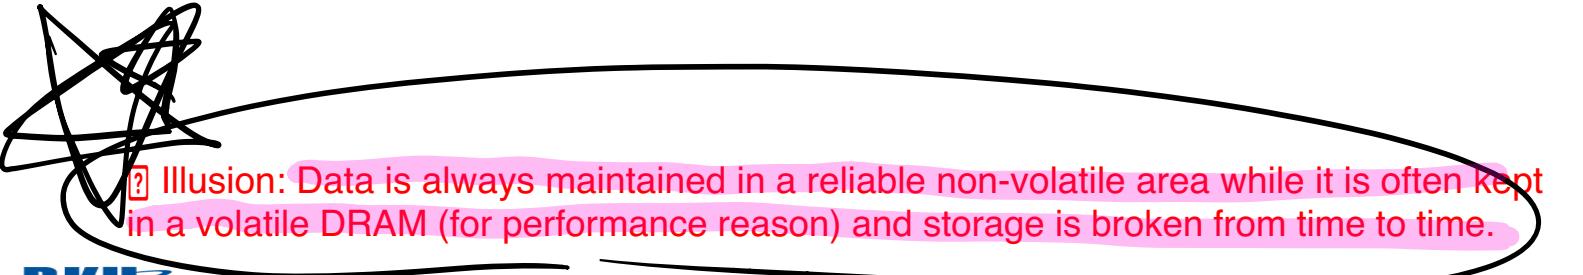
Be output
fd[0] fd[1] fd[2]
Be input *Be error.*

Figure 2.6: A Program That Does I/O (io.c)

2.4 Persistence

Issues for Persistence

- ✓ How to access a file?  `open()`, `read()`, `write()`, ...
- ✓ How to manage a file?  inode, FAT, ...
- ✓ How to manipulate a directory?
- ✓ How to design a file system?  UFS, LFS, Ext2/3/4, FAT, F2FS, NFS, AFS, ...
- ✓ How to find a data in a disk?
- ✓ How to improve performance in a file system?  cache, delayed write, ...
- ✓ How to handle a fault in a file system?  **journaling**, **copy-on-write**
- ✓ What is a role of a disk device **driver**?
- ✓ What are the internals of a disk and SSD?
- ✓ What is the RAID?



 **Illusion:** Data is always maintained in a reliable non-volatile area while it is often kept in a volatile DRAM (for performance reason) and storage is broken from time to time.

2.5 Design Goals

Abstraction

- ✓ Focusing on relevant issues only while hiding details
E.g. Car, File system, Make a program without thinking of logic gates
- ✓ “**Abstraction is fundamental to everything we do in computer science**” by Remzi

Performance

- ✓ Minimize the overhead of the OS (both time and space)

Protection

- ✓ Isolate processes from one another
- ✓ Access control, security, ...

Reliability

- ✓ Fault-tolerant

Others

- ✓ Depend on the area where OS is employed
- ✓ Real time, Energy-efficiency, Mobility, Load balancing, Autonomous, ...

2.5 Design Goals

Separation of Policy and Mechanism

✓ Policy: Which (or What) to do?

e.g.) Which process should run next?

✓ Mechanism: How to do?

e.g.) Multiple processes are managed by a scheduling queue or RB-tree

↑↑↑↑↑



That's a policy

That's a mechanism



That's a different
type of mechanism

(Source: Security Principles and Policies CS 236 On-Line MS Program Networks and Systems Security, Peter Reiher, Spring, 2008)

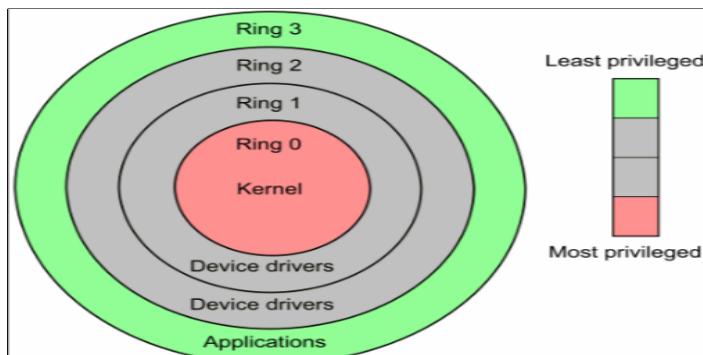
Early Operating Systems: Just libraries

- ✓ Commonly-used functions such as low-level I/Os (e.g. MS-DOS)
- ✓ **Batch processing** 
a number of jobs were set up and then run all together (Not interactive)

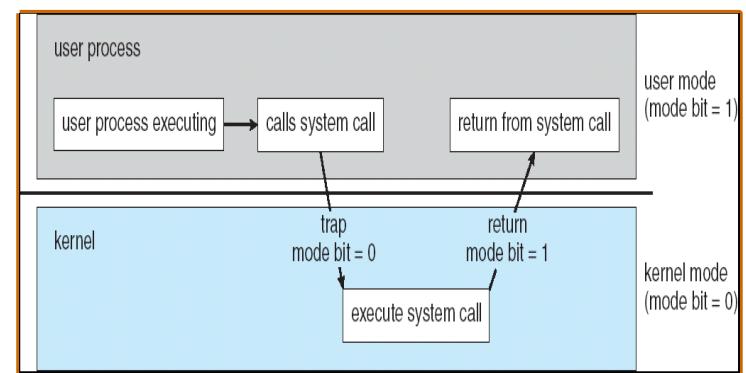
Beyond Libraries: Protection

- ✓ Require OS to be treated differently than user applications
- ✓ Separation user/kernel mode, system call
- ✓ Use trap (special hardware instruction) to switch into the kernel mode

Transfer control to a pre-specific trap handler (system_call handler)



(Source: Google Image)



(Source: A. Silberschatz, "Operating system Concept")

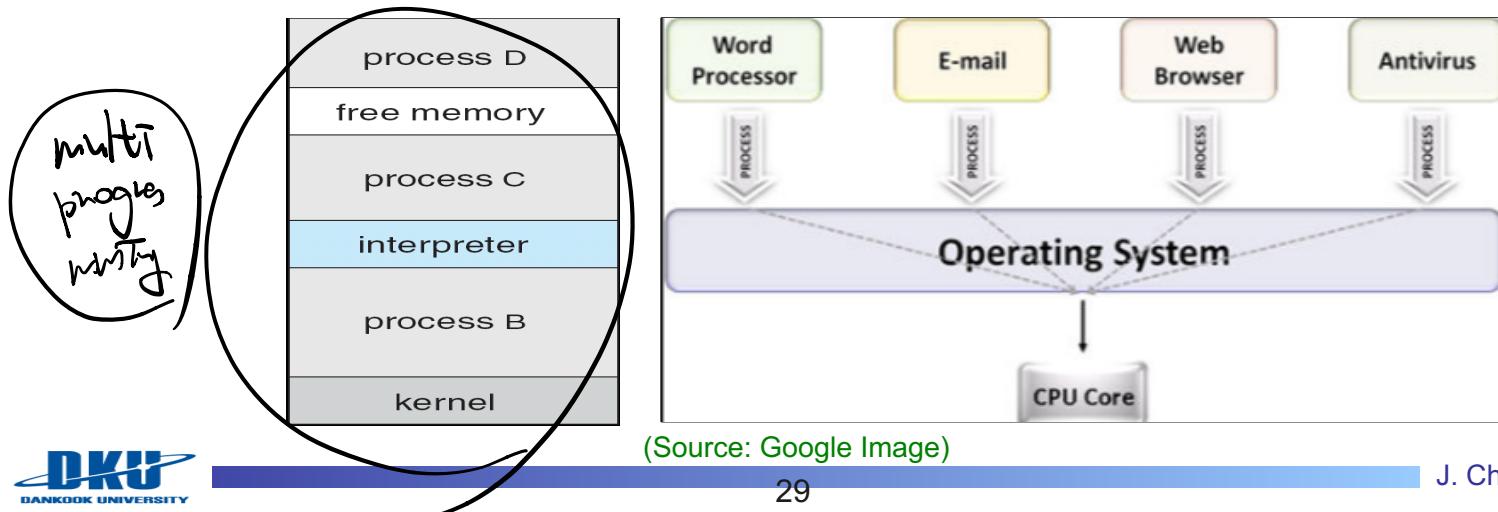
J. Choi, DKU

2.6 Some History

→ ~~历史~~ ~~发展~~ ~~技术~~

The Era of Multiprogramming (c.f. multitasking)

- ✓ Definition: OS load a number of applications into memory and switch them rapidly
- ✓ Reason: Advanced hardware → Want to utilize machine resources better → Multiple users share a system (workstation, minicomputer) → multiprogramming
- ✓ Especially important due to the slow I/O devices → while doing I/O, switch CPU to another process → enhancing CPU utilization
- ✓ Memory protection and concurrency become quite important → **UNIX**



2.6 Some History

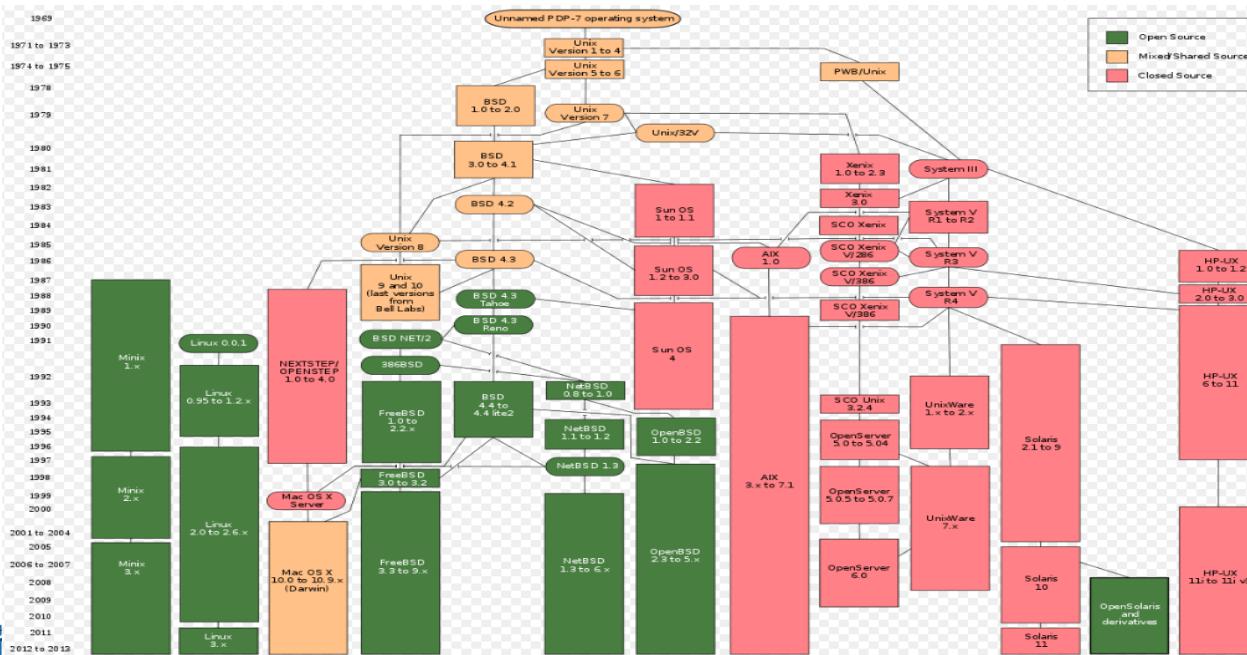
The Era of Multiprogramming (c.f. multitasking)

✓ UNIX

By Ken Thompson and Dennis Ritchie (Bell Labs), Influenced by Multics

C language based, excellent features such as shell, pipe, inode, small, everything is a file, ...

Influence OSes such as BSD, SUNOS, AIX, HPUX, Nextstep and Linux



2.6 Some History

The Modern Era

- ✓ PC
MS Windows, Mac OS X, Linux, ...
- ✓ Smartphone
Android, iOS, Windows Mobile, ...
- ✓ IoT
What is the next?



2.7 Summary

OS

- ✓ Resource manager (Efficiency)
- ✓ Make systems easy to use (Convenience)

Cover in this book

- ✓ Virtualization, Concurrency, Persistence

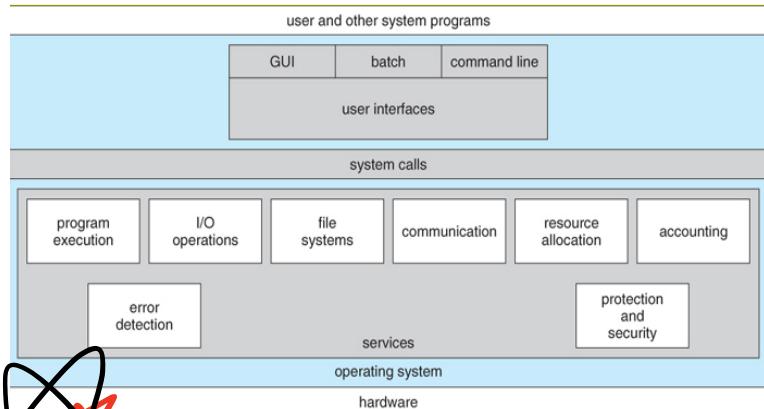
Not being covered

- ✓ Network, Security, Graphics
- ✓ There are several excellent courses for them

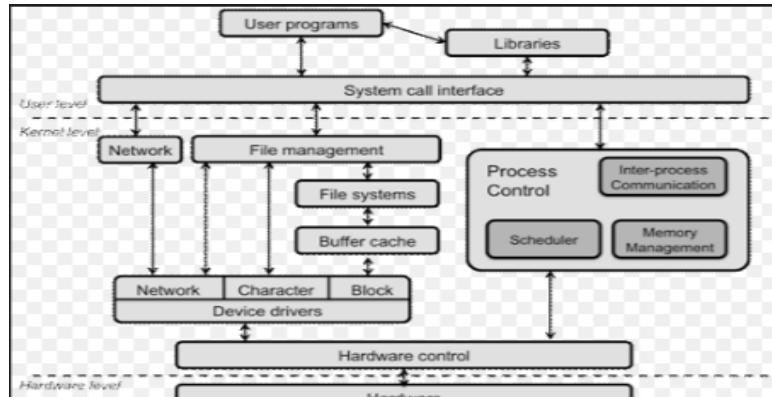
- Homework 1: summarize the chap 2 of the OSTEP
 - Requirement: 1) personal, 2) 3~5 pages for summary, 3) 1 page for the goal you want to study
 - Due: until the same day of the next week
 - Bonus: Snapshot of the results of example programs in a Linux system
- Any questions? Send an email to me: choijm@dankook.ac.kr
- Simple question to take attendance: Explain the differences between interrupt and trap which was discussed in page 28 (until 6 PM, March 24th)

Appendix

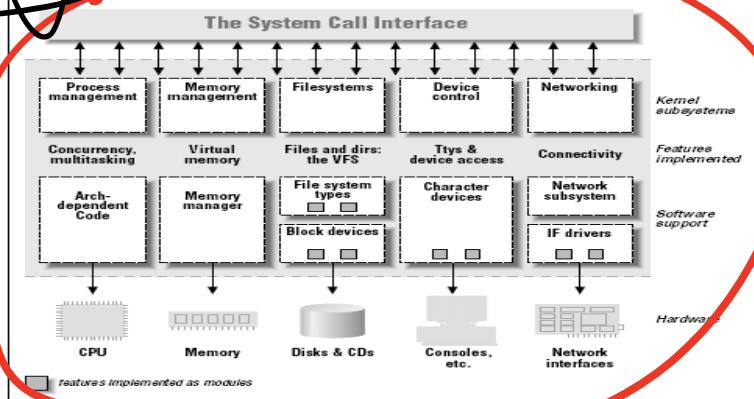
OS structure in General



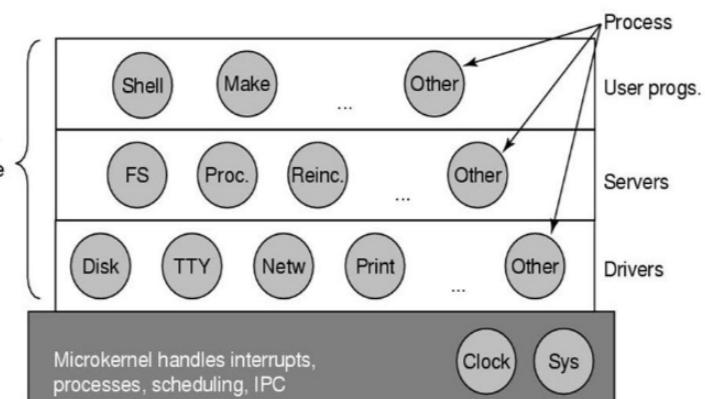
(Source: Operating System Concepts)



(Source: <https://www.cs.rutgers.edu/~pxk/416/notes/03-concepts.html>)



(Source: Linux Device Driver)



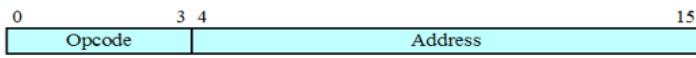
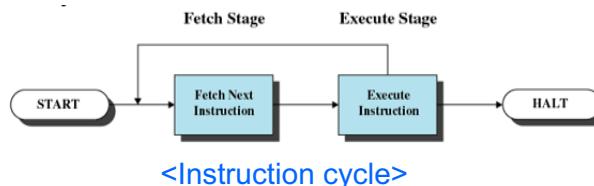
(Source: Modern Operating System)

Appendix

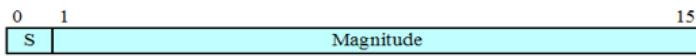
What happens when a program runs?

- ✓ Details: execute instructions

Fetch and Execute



(a) Instruction format



(b) Integer format

Program Counter (PC) = Address of instruction
 Instruction Register (IR) = Instruction being executed
 Accumulator (AC) = Temporary storage

(c) Internal CPU registers

0001 = Load AC from Memory
 0010 = Store AC to Memory
 0101 = Add to AC from Memory

(d) Partial list of opcodes

<Hypothetical machine>

(Source: W. Stallings, "Operating Systems: Internals and Design Principles")

