

Lecture Note 8: Memory Management

May 27, 2020
Jongmoo Choi

Dept. of Software
Dankook University

<http://embedded.dankook.ac.kr/~choijm>

(This slide is made by Jongmoo Choi. Please let him know when you want to distribute this slide)

Contents

From Chap 12~17 of the OSTEP

Chap 12. A Dialogue on Memory Virtualization

Chap 13. The Abstraction: Address Space

Chap 14. Interlude: Memory API

- ✓ malloc(), free(), brk(), mmap(), ...

Chap 15. Mechanism: Address Translation

- ✓ Base & Limit (Bound), Dynamic Relocation

Chap 16. Segmentation

- ✓ Generalization, Sharing, Protection

Chap 17. Free-Space Management

- ✓ Fragmentation, Splitting and Coalescing
- ✓ Strategies: Best fit, First fit, Worst fit, ...
- ✓ Segregated list, Buddy algorithm, ...

Chap 12. Dialogue

Memory virtualization

Student: So, are we done with virtualization?

Professor: No!

Student: Hey, no reason to get so excited; I was just asking a question. Students are supposed to do that, right?

Professor: Well, professors do always say that, but really they mean this: ask questions, if they are good questions, and you have actually put a little thought into them.

Student: Well, that sure takes the wind out of my sails.

Professor: Mission accomplished. In any case, we are not nearly done with virtualization! Rather, you have just seen how to virtualize the CPU, but really there is a big monster waiting in the closet: memory. Virtualizing memory is complicated and requires us to understand many more intricate details about how the hardware and OS interact.

Student: That sounds cool. Why is it so hard?

Professor: Well, there are a lot of details, and you have to keep them straight in your head to really develop a mental model of what is going on. We'll start simple, with very basic techniques like base/bounds, and slowly add complexity to tackle new challenges, including fun topics like TLBs and multi-level page tables. Eventually, we'll be able to describe the workings of a fully-functional modern virtual memory manager.

Student: Neat! Any tips for the poor student, inundated with all of this information and generally sleep-deprived?

Professor: For the sleep deprivation, that's easy: sleep more (and party less). For understanding virtual memory, start with this: every address generated by a user program is a virtual address. The OS is just providing an illusion to each process, specifically that it has its own large and private memory; with some hardware help, the OS will turn these pretend virtual addresses into real physical addresses, and thus be able to locate the desired information.

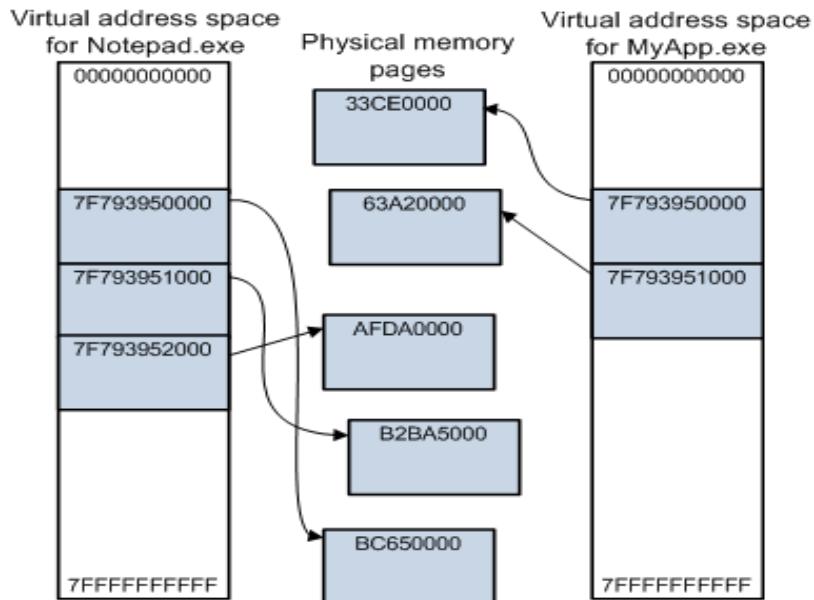
Chap 13. The abstraction: address space

Early system

Multiprogramming and Time sharing

Address space

Goals



(Source: <https://msdn.microsoft.com/en-us/windows/hardware/drivers/gettingstarted/virtual-address-spaces>)

13.1 Early Systems

Use physical memory directly

- ✓ OS and current program single programming system
- ✓ No (limited) protection
- ✓ Larger program than physical memory **Overlay**

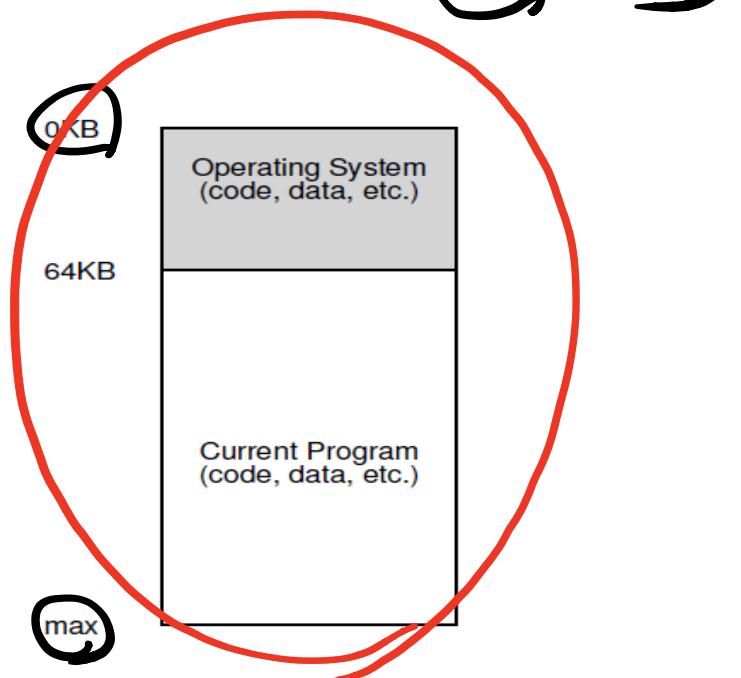


Figure 13.1: Operating Systems: The Early Days

13.2 Multiprogramming and Time sharing

Computer becomes bigger



Multiprogramming: multiple processes are ready to run

Time sharing: switch CPUs among ready processes

✓ Issues

Protection becomes a critical issue

How to find suitable free space

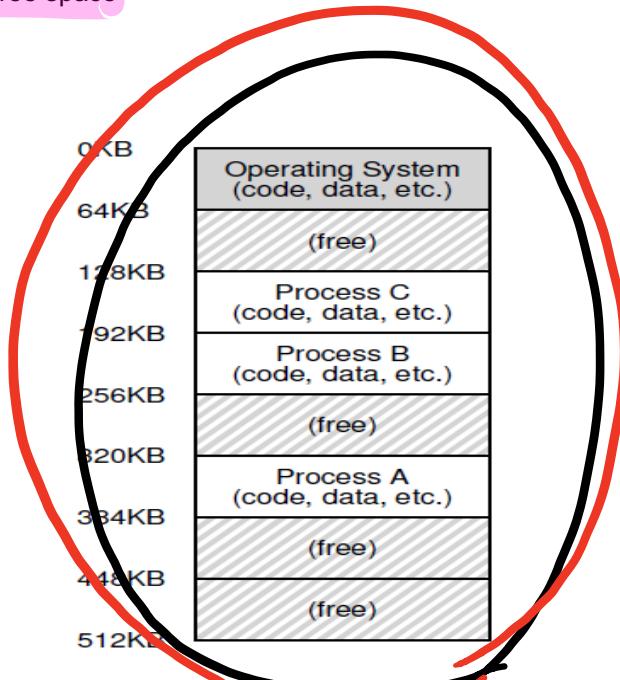


Figure 13.2: Three Processes: Sharing Memory

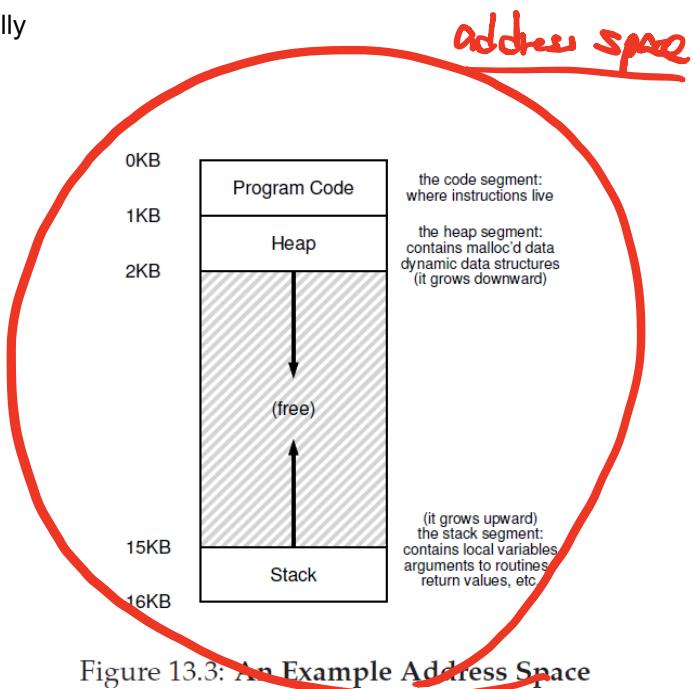
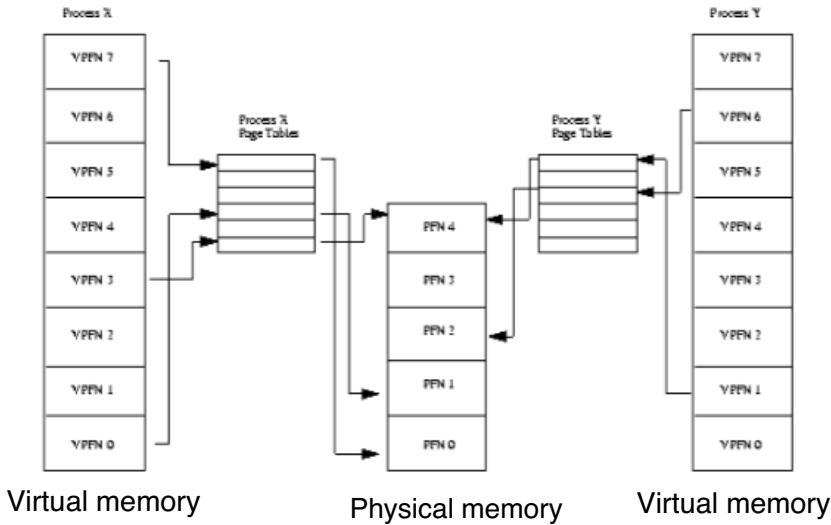
13.3 Address space

Abstraction

- ✓ A process has an illusion that it uses exclusively all memory even though it is shared by multiple processes **virtual memory**
- ✓ Well defined layout **address space**

Code (instruction), Data (statically-initialized variables), Stack (function call chain and local variables), Heap (dynamically allocated)

Code is located at virtual address 0x0, but not physically



13.4 Goals

- Transparency (easy to use)

- ✓ Programmer: no need to aware the memory size or available space

- Efficiency

- ✓ Both in terms of time and space (not slow and not requires much additional overhead)  Various HW support (e.g. TLB)

- Protection (isolation)

- ✓ Protect processes from one another

 Note: **every address you see is virtual**

```
#include <stdio.h>
#include <stdlib.h>
int main(int argc, char *argv[]) {
    printf("location of code : %p\n", (void *) main);
    printf("location of heap : %p\n", (void *) malloc(1));
    int x = 3;
    printf("location of stack : %p\n", (void *) &x);
    return x;
}
```

When run on a 64-bit Mac OS X machine, we get the following output:

```
location of code : 0x1095afe50
location of heap : 0x1096008c0
location of stack : 0x7fff691aea64
```

Chap 14. Interlude: Memory API

Types of Memory

The malloc() call

The free() call

Common errors

Underlying OS Support

Other Calls



`# free()`



`# malloc([])`

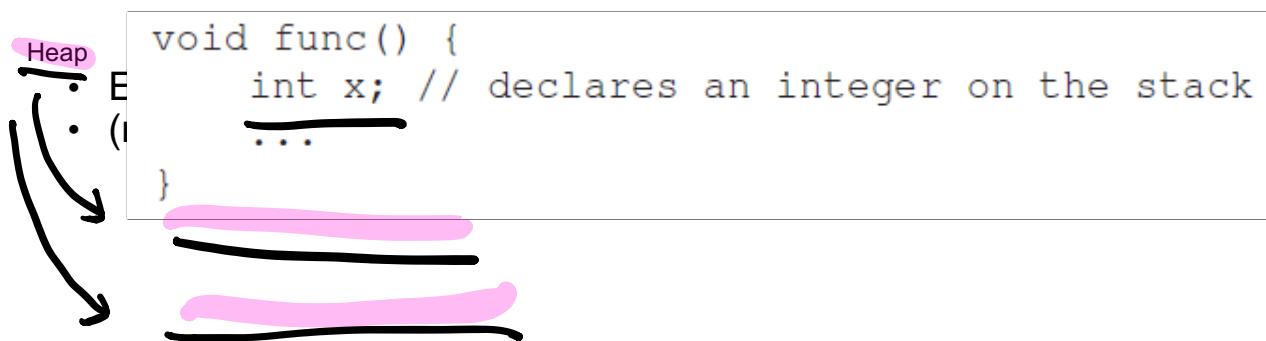


14.1 Types of Memory

Two types of memory

- ✓ Static: Code (also called as text), Data
- ✓ Dynamic: Heap, Stack

Stack
Implicitly by the compiler (hence sometimes called automatic memory)
Short-lived memory



```
void func() {
    int *x = (int *) malloc(sizeof(int));
    ...
}
```

14.2/3 The malloc()/free() call

The malloc() call

- ✓ Input: memory size (how many bytes you need)
- ✓ Output: pointer to the newly-allocated space (or NULL if it fails)
- ✓ Use well-defined macros or routines, instead of number as input

```
✓ malloc(sizeof(int));  
    malloc(strlen(s) + 1);
```

The free() call

- ✓ Input: pointer (size is not specified, meaning that it is managed by the library)

```
int *x = malloc(10 * sizeof(int));  
...  
free(x);
```

14.4 Common errors

Common errors

✖ Forgetting to allocate memory

```
char *src = "hello";
char *dst;           // oops! unallocated
strcpy(dst, src); // segfault and die
```

We frequently meet the [segmentation fault](#). Hence ↗

```
char *src = "hello";
char *dst = (char *) malloc(strlen(src) + 1);
strcpy(dst, src); // work properly
```

When you run this code, it will likely lead to a [segmentation fault](#)³, which is a fancy term for **YOU DID SOMETHING WRONG WITH MEMORY YOU FOOLISH PROGRAMMER AND I AM ANGRY.**

☞ [Make use of a debugger \(e.g. gdb\)](#)

14.4 Common errors

1가지

☞ → Past Ab

Common errors

- ✓ Not allocating enough Memory

```
char *src = "hello";
char *dst = (char *) malloc(strlen(src)); // too small! abilities.
strcpy(dst, src); // work properly
```

- ✓ F
Heap has data of unknown value.
- ✓ Forgetting to free memory

Memory leak

Some languages support the garbage collection mechanism that manages memory automatically without requiring explicit free() by programmers [?] but if you still have a reference, the collector will never free it (still problem)

14.4 Common errors

Common errors

- ✗ ✓ Freeing memory before you are done with it

Dangling pointer

Subsequent use can crash the program and even system



- ✓ Freeing memory repeatedly

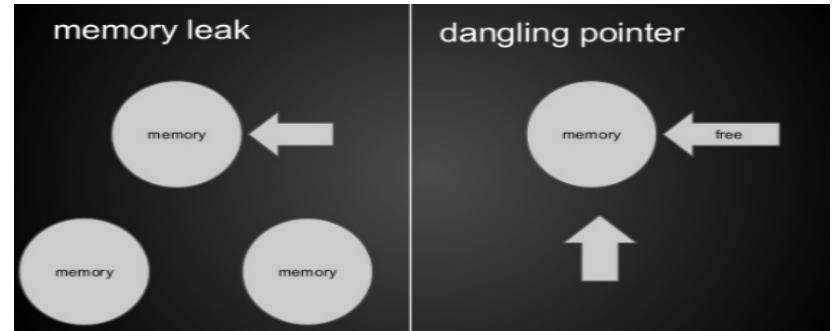
Double free

- ✓ Calling free() incorrectly

Invalid free

Tools for solving memory-related problems

- ✓ Purify
- ✓ Valgrind
- ✓ ...



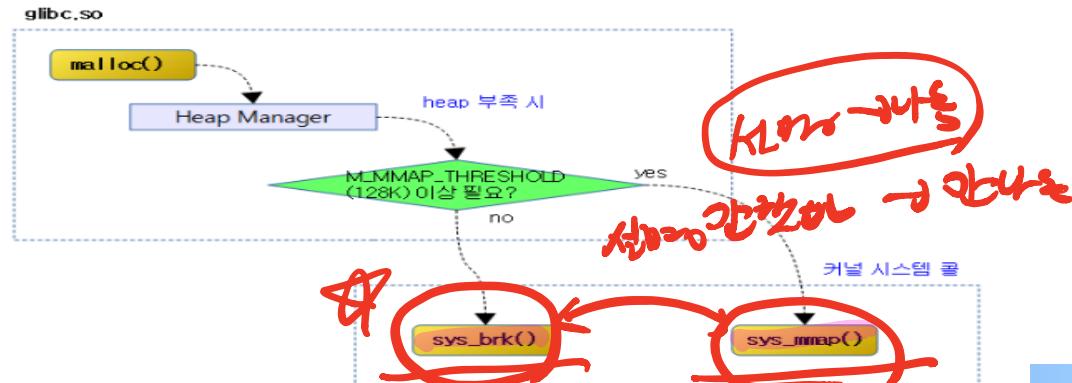
14.5/6 Underlying OS Support/Other Calls (Optional)

Underlying OS Support

- ✓ `malloc()/free()`  library
- ✓ It internally allocates several pages using the `sys_brk()` or `sys_mmap()` system call and manages them to serve the `malloc()` and `free()` request
- ✓ If its space becomes too small, it requests more pages to OS again using the ~~`sys_brk()`~~ or `sys_mmap()`  system call

Other Calls

- (`calloc()`: allocate and zero space
- (`realloc()`: allocate a new larger region, copy the old region into it and returns the pointer of the new region



Chap. 15 Mechanism: Address Translation

CPU virtualization

- ✓ Limited Direct Execution

Direct execution: process run independently for the most time (efficiency)
Limited: OS get involved (control)

- ✓ Two mechanisms

Restricted operations (e.g. system call): user mode $\xrightarrow{?}$ kernel mode (OS control)
Timer interrupt: user mode $\xrightarrow{?}$ kernel mode (OS control), do periodic jobs such as scheduling and context switch

Memory virtualization

~~Address Translation~~

Address space: virtual memory (using virtual address)

During execution: physical memory (using physical address which is translated from virtual address)

- ✓ Again, we will pursue both ~~efficiency and control~~

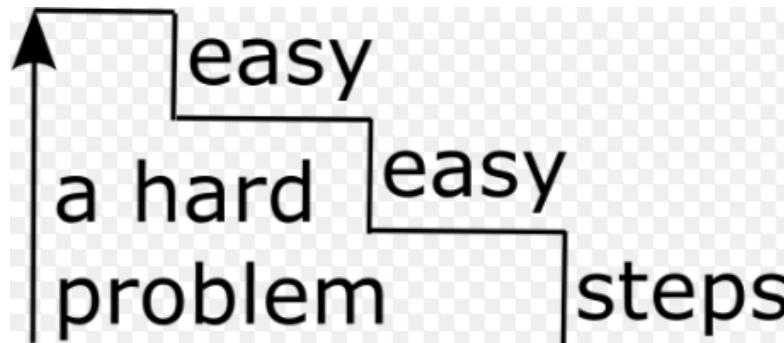
Efficiency: small overhead $\xrightarrow{?}$ hardware-based address translation

Control: OS ensures that no processes is allowed to access any memory but its own $\xrightarrow{?}$ OS memory management

15.1 Assumption

Three assumptions for easy explanation

- ✓ Address space must be placed contiguously in physical memory
 - ✓ The size of address is not too big (less than physical memory)
 - ✓ Each address space is the same size
- ?
- We will relax these as we go.



15.2 An Example

A program

- ✓ High-level viewpoint

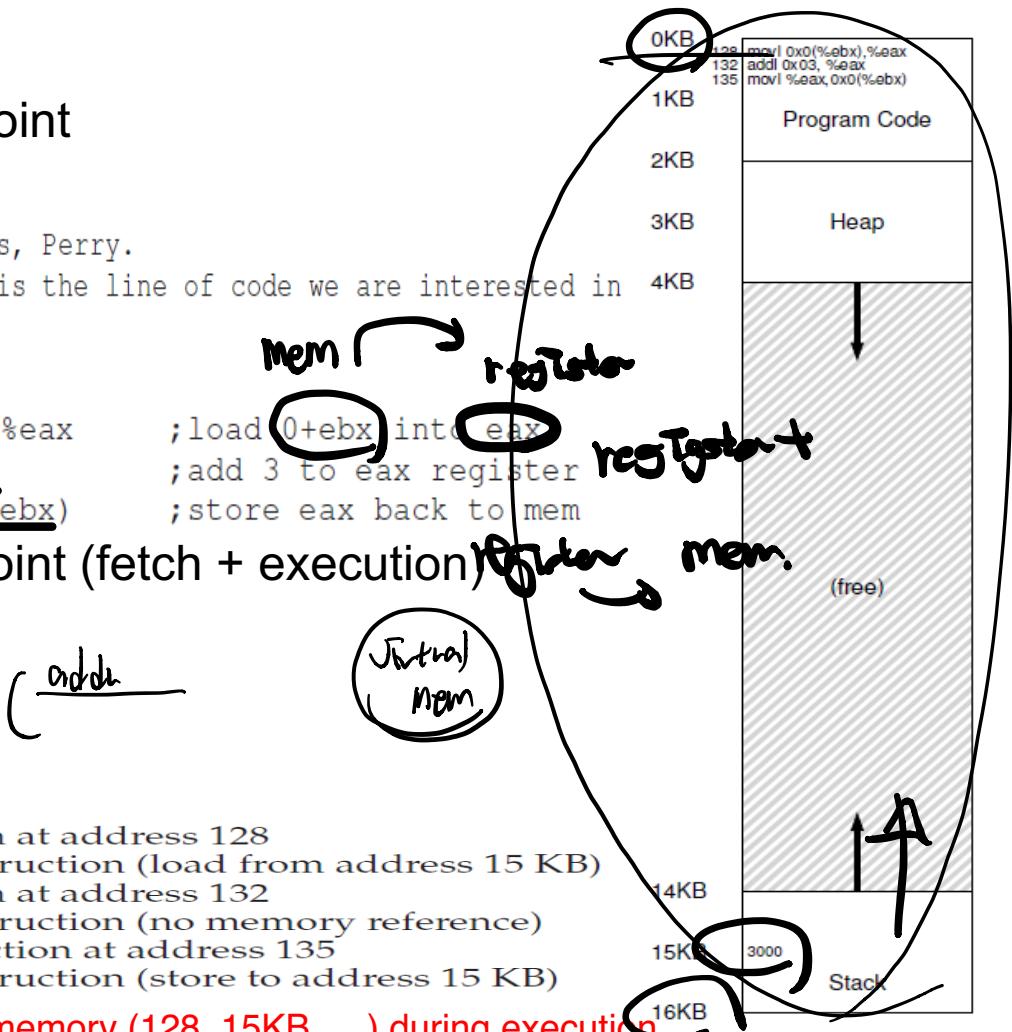
```
void func() {  
    int x = 3000; // thanks, Perry.  
    x = x + 3; // this is the line of code we are interested in
```

```
128: movl 0x0(%ebx), %eax ;load 0+ebx into eax  
132: addl $0x03, %eax ;add 3 to eax register  
135: movl %eax, 0x0(%ebx) ;store eax back to mem
```

- ✓ Execution viewpoint (fetch + execution)

- Fetch instruction at address 128
- Execute this instruction (load from address 15 KB)
- Fetch instruction at address 132
- Execute this instruction (no memory reference)
- Fetch the instruction at address 135
- Execute this instruction (store to address 15 KB)

?? Need to access memory (128, 15KB, ...) during execution



15.2 An Example

Focusing on memory

✓ Address space (virtual memory)

Starts at address 0

Grows to maximum of 16 KB

1
2
3

✓ Physical memory

Can place any free space, not necessarily at address 0 **relocation**

✓ Address translation

Assume that the process is located at 32KB~48KB in physical memory

Then the virtual address 0 needs to be translated into the physical address 32KB **address translation**

✓ Other slots (e.g. 16~32KB) are not used

free space management

✓ OS also locates in physical memory

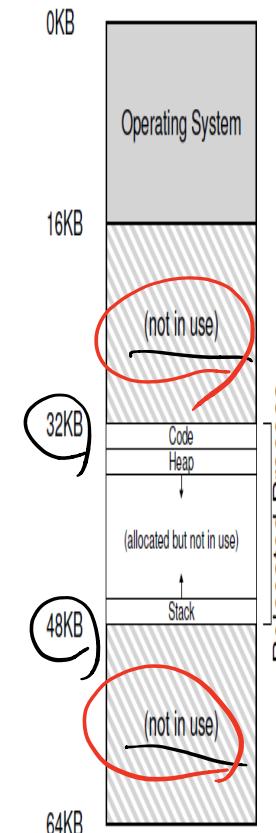
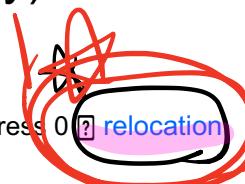


Figure 15.2: Physical Memory with a Single Relocated Process

15.3 Dynamic (Hardware-based) Relocation

Integration of Virtual and Physical memory

- ✓ Virtual memory: 0~16KB vs Physical Memory: 0~64KB

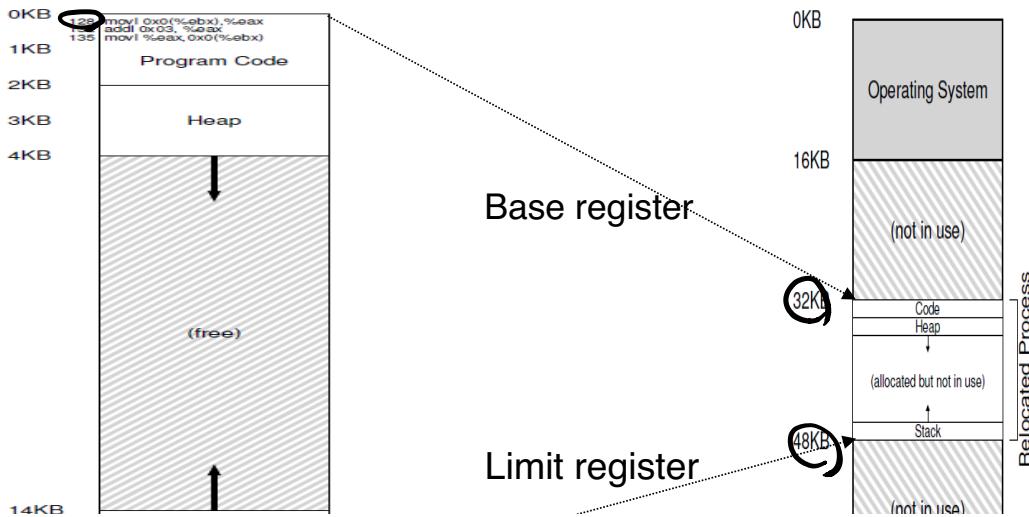
Being loaded into 32KB 48KB

- ✗ **Address translation:** virtual address \rightarrow physical address

First instruction 128 \rightarrow 32KB + 128 ($32768 + 128 = 32896$)

Variable x: 15KB \rightarrow 32KB + 15KB = 47KB

In general: base address + offset (instruction or variable's address)



- Simple question to take attendance: Imagine a process with an address space of size 4KB has been loaded at physical address 16KB. 1) What are the values of base and limit registers? 2) What are the corresponding physical addresses for virtual addresses of 1) 1KB, 2) 3000 and 3) 4400. (hint: see page 8 of chap. 15 in OSTEP) (until 6 PM, June 3rd)

15.3 Dynamic (Hardware-based) Relocation

Summary of address translation (and relocation)

- ✓ Virtual memory: per process (exclusive), start at 0x0 (size: 16KB)
- ✓ Physical memory: shared by processes, start at any address (different among processes)
- ✓ Three main components: Compiler, OS and Hardware (Architecture)
 - (A program is **compiled** as if it is loaded at address 0 (virtual memory).
 - (The program is **loaded** any space in physical memory, while setting base and limit registers appropriately **relocatable**
 - (An address requested by CPU is **translated** into a physical address while running (and protected)

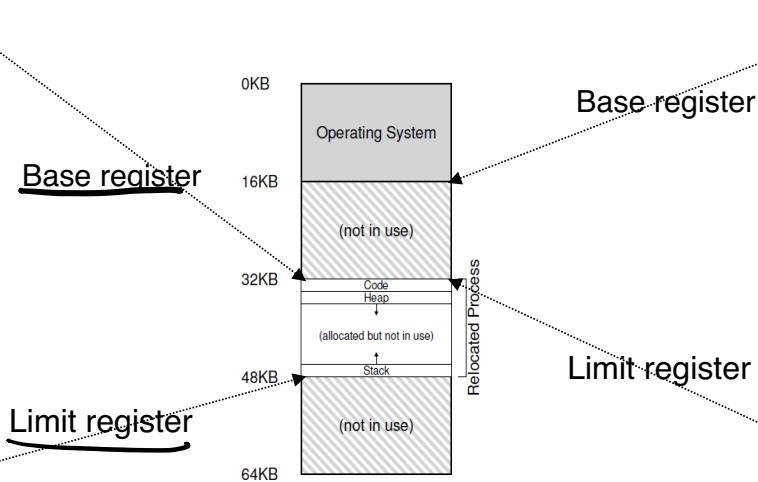
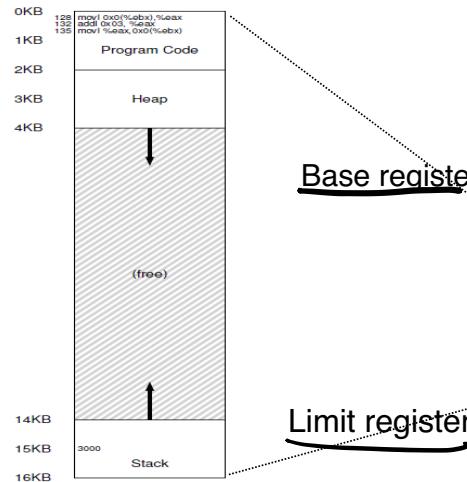


Figure 15.1: A Process And Its Address Space

Figure 15.2: Physical Memory with a Single Relocated Process

Virtual memory (for process A)

Physical memory

Figure 15.1: A Process And Its Address Space

Figure 15.2: Physical Memory with a Single Relocated Process

Virtual memory (for process B)

15.3 Dynamic (Hardware-based) Relocation

Summary of address translation (and relocation)

- ✓ How to translate? Using two hardware registers

Base register: start address (30004 in this example)

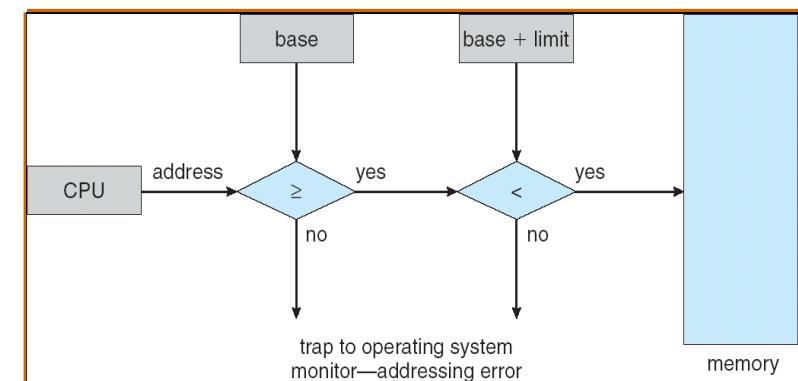
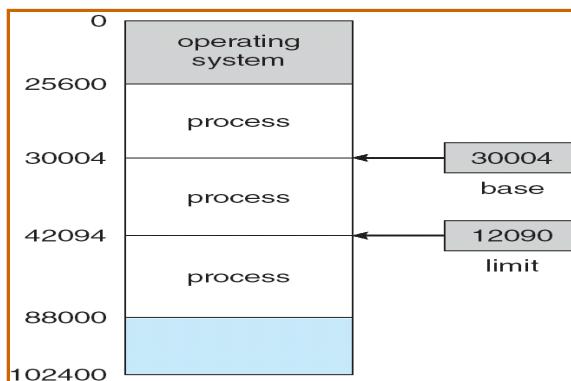
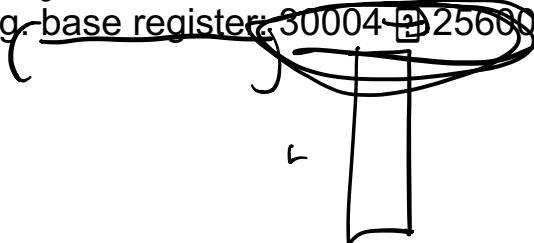
- physical address = base register + virtual address
- E.g. virtual address = 10 physical address = 30014

Limit register (Bound register): upper bound (or size, 12090 in this example)

- E.g. virtual address = 13000 physical address = segmentation fault

Base/Limit registers are switched at each context switch time

- E.g. base register: 30004 → 25600



(Source: A. Silberschatz, "Operating system Concept")

15.4 Hardware Support: A Summary

MMU (Memory management unit)

- Part of CPU that helps with address translation
- ✓ E.g.) Base/limit registers, Segmentation related registers, Paging related registers, ~~TLB (Translation Lookaside Buffer)~~ + circuitry

Summary of HW support for Dynamic relocation

Hardware Requirements	Notes
Privileged mode	Needed to prevent user-mode processes from executing privileged operations
Base/bounds registers	Need pair of registers per CPU to support address translation and bounds checks
Ability to translate virtual addresses and check if within bounds	Circuitry to do translations and check limits; in this case, quite simple
Privileged instruction(s) to update base/bounds	OS must be able to set these values before letting a user program run
Privileged instruction(s) to register exception handlers	OS must be able to tell hardware what code to run if exception occurs
Ability to raise exceptions	When processes try to access privileged instructions or out-of-bounds memory

Figure 15.3: Dynamic Relocation: Hardware Requirements

15.5 Operating Systems Issues

OS responsibilities

~~Memory management~~

Allocation for new processes, free list manipulation, ...
Reclaim the space of terminated processes

~~Base/limit registers management during Context switch~~

Save/restore base/limit registers into/from PCB (MMU)
Process **relocation** if necessary

~~Exception handling~~

Handlers + Table (e.g. segmentation fault handler + IVT)

OS Requirements

Memory management

Notes

*Need to allocate memory for new processes;
Reclaim memory from terminated processes;
Generally manage memory via free list*

Base/bounds management

Must set base/bounds properly upon context switch

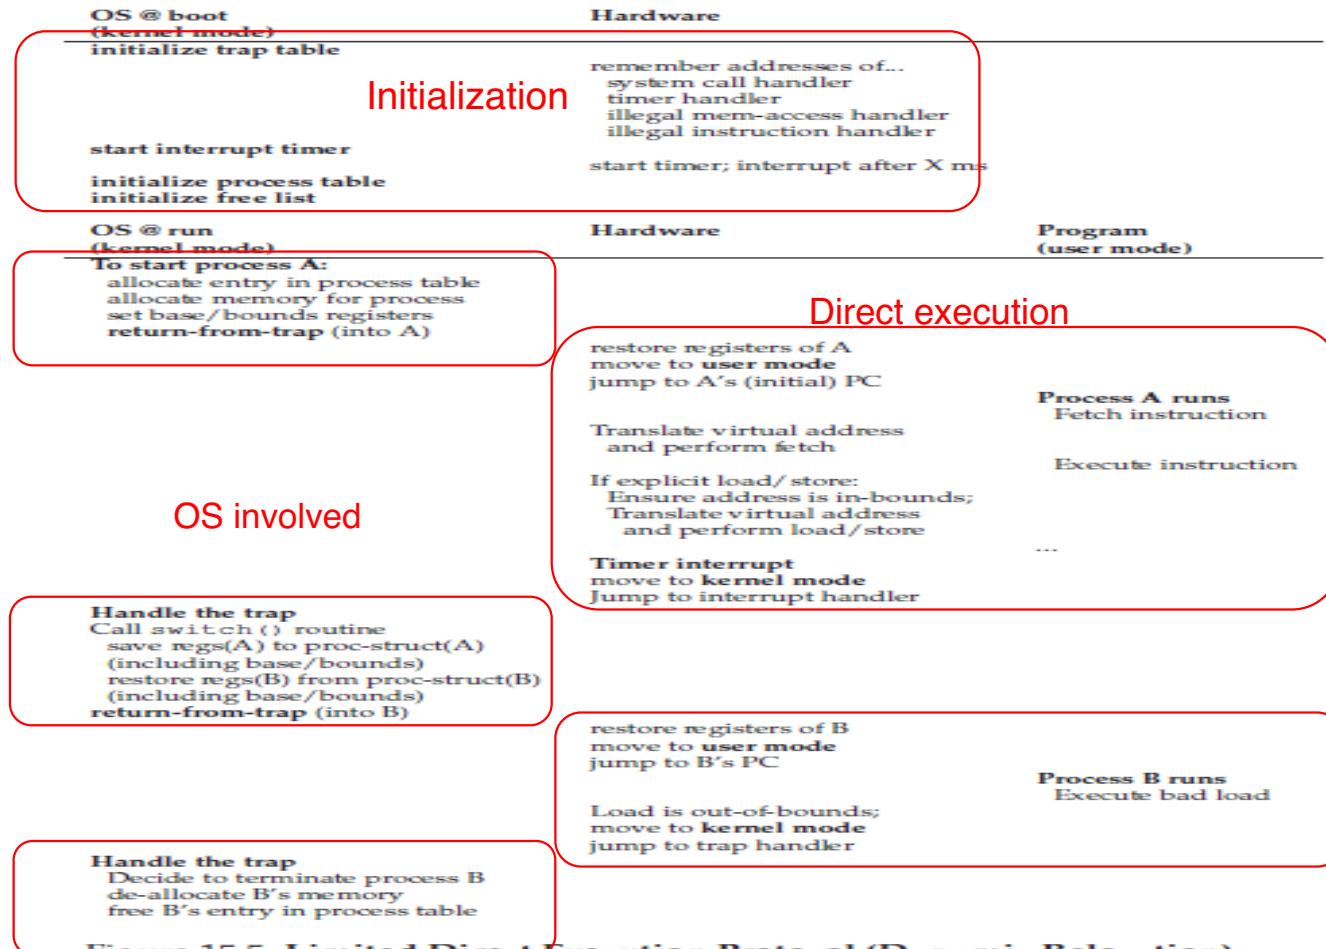
Exception handling

*Code to run when exceptions arise;
likely action is to terminate offending process*

Figure 15.4: Dynamic Relocation: Operating System Responsibilities

15.5 Operating Systems Issues

Global view



15.6 Summary

Memory virtualization

Address translation

OS: memory allocation/free, base/limit initialize, exception control (infrequent event)

HW: virtual to physical at every execution (frequent event, MMU)

Support transparency: users have no idea where their processes are

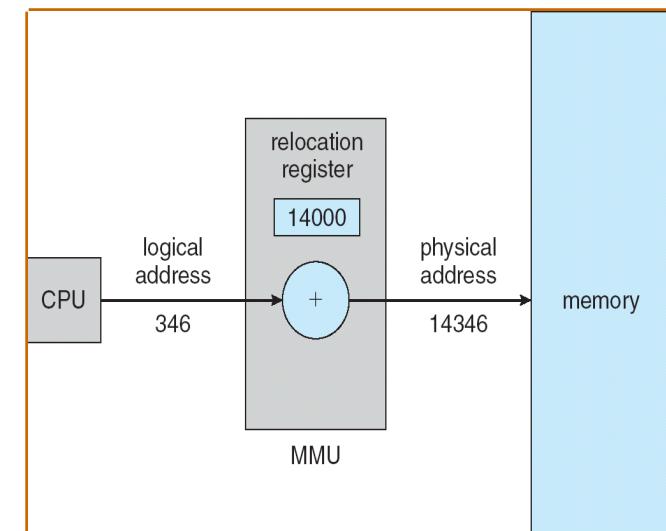
Mechanisms

Contiguous allocation

- 1) Base and limit registers
- Pros: Simple and Offer protection
- Cons: Internal fragmentation

Non-contiguous allocation

- 2) Segmentation: Variable size
- 3) Paging: Fixed size

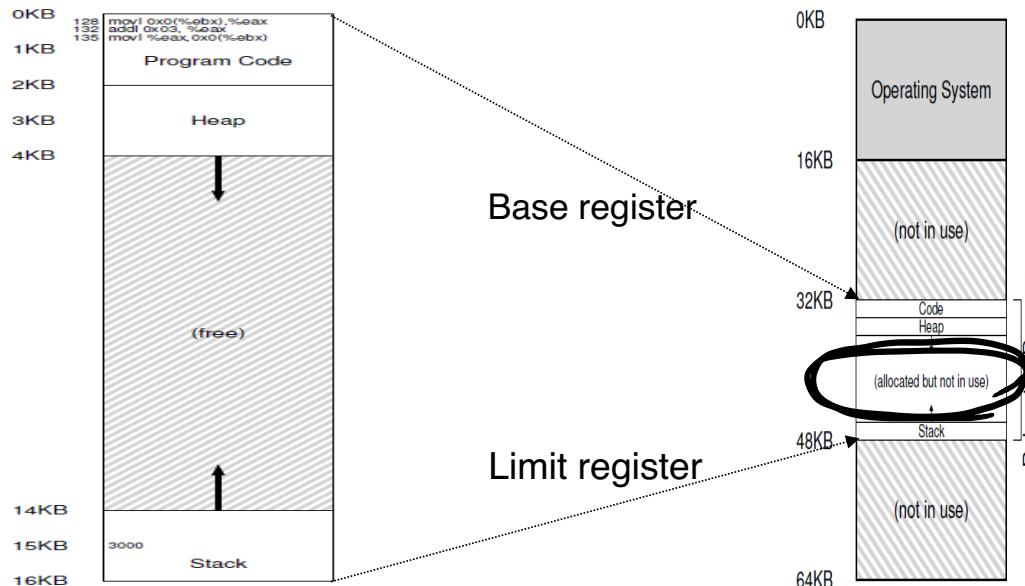


(Source: A. Silberschatz, "Operating system Concept")

Chap. 16 Segmentation

Issues of the base/limit register based dynamic relocation

- ✓ A big chunk of “free” space in the middle of address space
Even though they are free, they are taking up physical memory
- ✓ Hard to run a program when the entire address space does not fit into an available space in physical memory



?] How large the free space between heap and stack in 32-bit CPU?

16.1 Segmentation: Generalized Base/Limits

Key idea

- ✓ Contiguous **?** Non-contiguous
- ✓ Segment: divide a program into multiple segments (each segment is a contiguous portion of the address space)
E.g.) code segment, data segment, stack segment, heap segment, ...
- ✓ **Support base/limit per segment**
OS places segments independently in physical memory

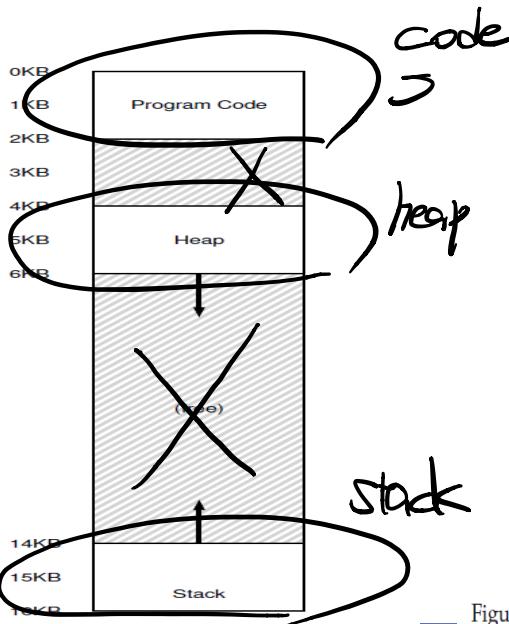


Figure 16.1: An Address Space (Again)

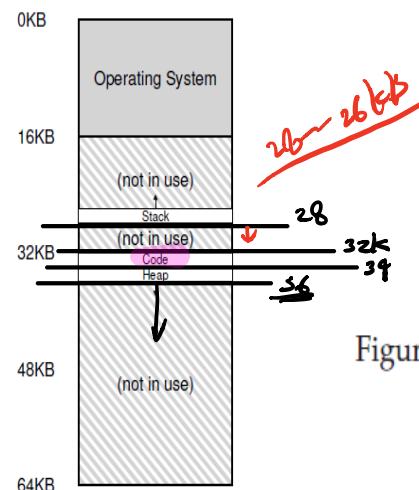


Figure 16.2: Placing Segments In Physical Memory

Segment	Base	Size
Code	32K	2K
Heap	34K	2K
Stack	28K	2K

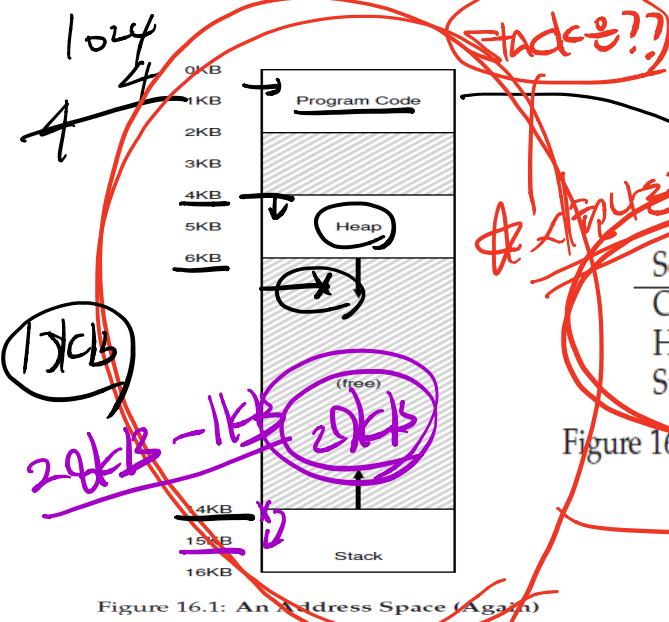
Figure 16.3: Segment Register Values

16.1 Segmentation: Generalized Base/Limits

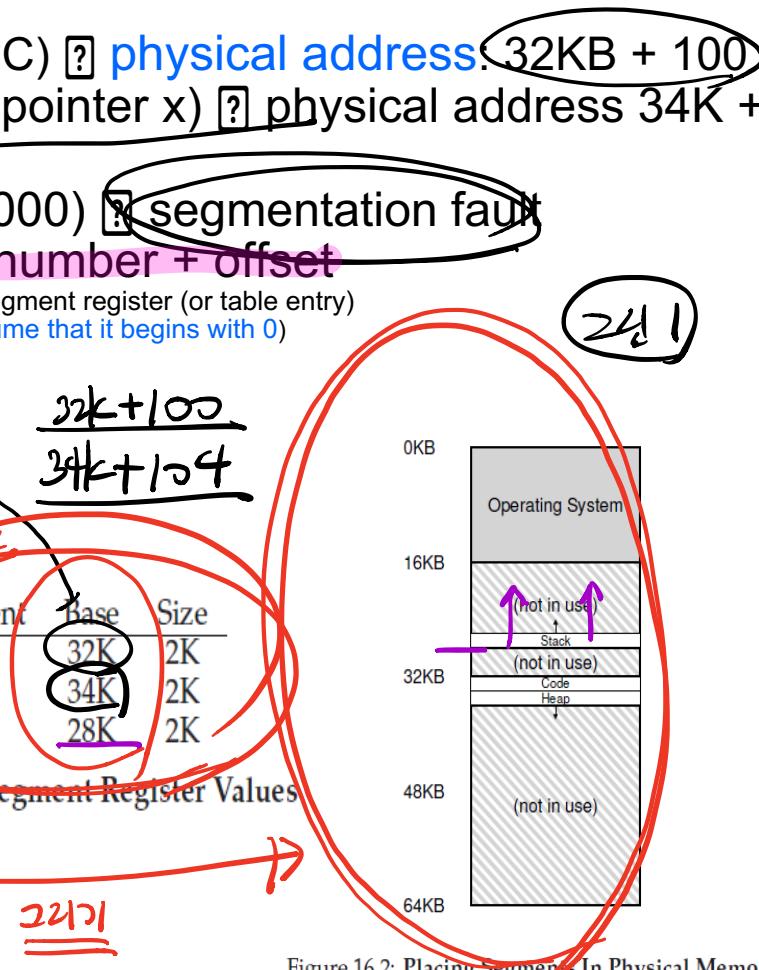
Address translation

- ✓ virtual address 100 (e.g. PC) ↗ physical address: 32KB + 100
 - ✓ virtual address 4200 (e.g. pointer x) ↗ physical address 34K + 104
 - ✓ virtual address 7000 (or 3000) ↗ segmentation fault
 - ✓ virtual address: segment number + offset

Segment number: choose appropriate segment register (or table entry)
Offset: location within the segment (assume that it begins with 0)



~~Figure 16.1: An Address Space (Again)~~



~~Figure 16.3: Segment Register Values~~

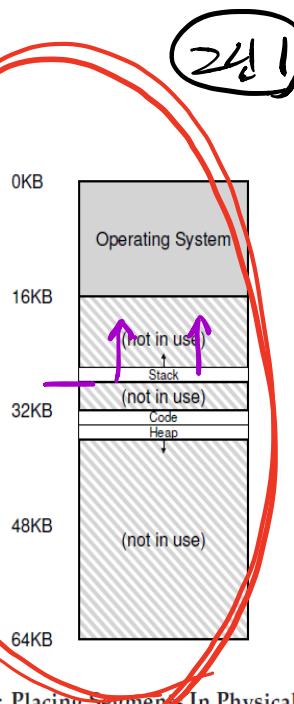


Figure 16.2: Placing Segments In Physical Memory

16.2 Which Segmentation Are We Referring To?

Segment encoding in virtual address

- ✓ Segment number part + offset part
- ✓ In the previous example

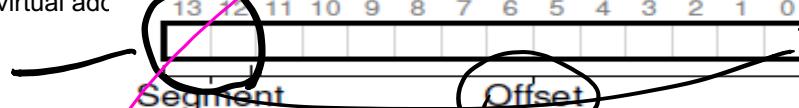
Address space size: $16KB = 2^{14}$ \Rightarrow 14 bit

Number of segment: $3 \Rightarrow 2$ bit

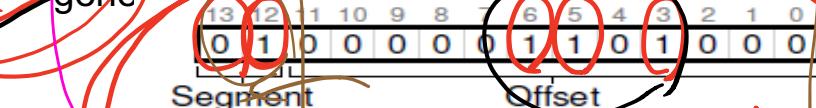
Number of offset: remaining 12 bit \Rightarrow maximum size of a segment: $4KB$

14bit

Segment: 00 \Rightarrow code, 01 \Rightarrow heap, 11 \Rightarrow stack
virtual addr



- Segment number: Used for searching its related base register
- Offset: If this offset is larger than the limit, trigger the segmentation fault. Otherwise, add offset with the value of the base register to generate the physical address ($1200 \Rightarrow 01$ (heap) + 1000)



heap

Stack

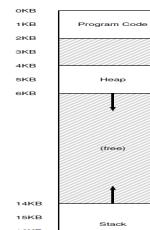


Figure 16.1: An Address Space (Again)

Segment	Base	Size
Code	32K	2K
Heap	34K	2K
Stack	28K	2K

Figure 16.3: Segment Register Values

16.2 Which Segmentation Are We Referring To?

Pseudo code for Segmentation

```
1 // get top 2 bits of 14 bit VA
2 Segment = (VirtualAddress & SEG_MASK) >> SEG_SHIFT
3 // now get offset
4 Offset = VirtualAddress & OFFSET MASK
5 if (Offset >= Bounds[Segment])
6   RaiseException(PROTECTION FAULT)
7 else
8   PhysAddr = Base[Segment] + Offset
9   Register = AccessMemory(PhysAddr)
```

~~PROTECTION FAULT~~

- Simple question to take attendance: 1) From the above pseudo code, what are the values of SEG_MASK, OFFSET_MASK and SEG_SHIFT? (assume that the size of virtual memory is 16KB and the size of each segment is 4KB, like in page 29), 2) Explain how to translate from virtual addresses 100 and 7000 into physical addresses (shown in page 29) using the bit presentation discussed in page 30 (until 6 PM, June 4th).



16.3 What About the Stack?

16.3.2

Stack issue

- ✓ It grows backward \square translation must proceed differently
Need extra HW support

In
sp
) ✓

Segment	Base	Size (max 4K)	Grows Positive?
Code ₀₀	32K	2K	1
Heap ₀₁	34K	3K	1
Stack ₁₁	28K	2K	0

Figure 16.4: Segment Registers (With Negative-Growth Support)

Virtual address: 15KB = 11 1100 0000 0000

- Segment number 11 \square stack
- Offset 1100 0000 0000 \square 3KB

Physical address: 28KB + (15KB - 16KB) or 28KB + (3KB - 4KB) \square 27KB



Figure 16.2: Placing Segments in Physical Memory

16.4/5 Support for Sharing/ Granularity

Benefit of segmentation

- ✓ Sharing among multiple processes
- ✓ Protection support

Segment	Base	Size (max 4K)	Grows Positive?	Protection
Code ₀₀	32K	2K	1	Read-Execute
Heap ₀₁	34K	3K	1	Read-Write
Stack ₁₁	28K	2K	0	Read-Write

Figure 16.5: Segment Register Values (with Protection)

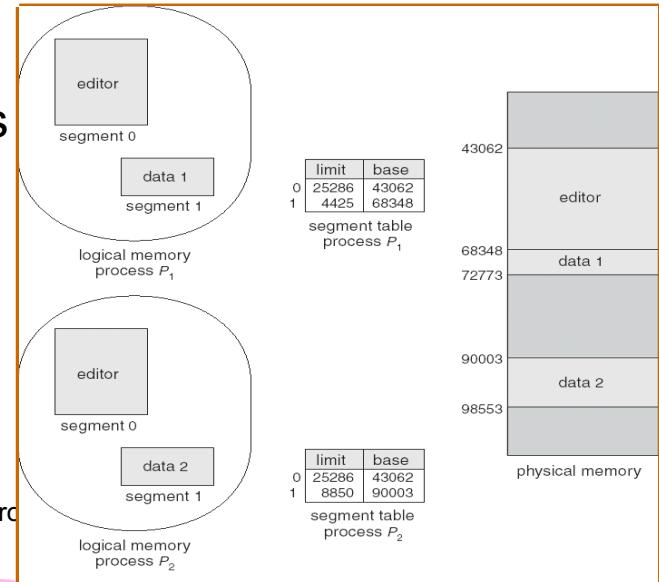
Relatively large size, small # of segments in a process (arc)

✓ Fine-grained

Relatively small size, large # of segments in a process

Make use of a table **segment table** for manipulating large # of segments.

(Source: A. Silberschatz, "Operating system Concept")



16.6 OS Support

For segmentation support

- Context switch: save/restore segment related registers
- Free space management

Try to reduce external fragmentation coalescing and compaction

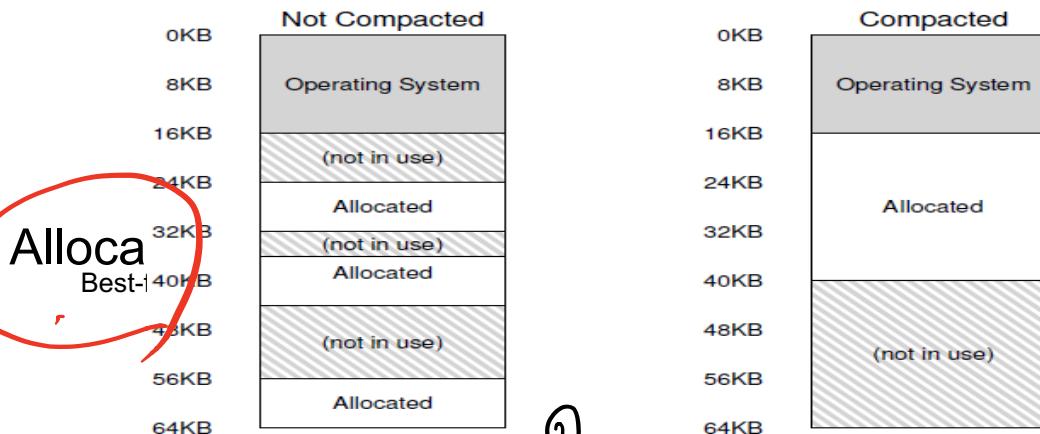


Figure 16.6: Non-compact and Compacted Memory

Compaction in memory: prepare for large free space vs Compaction in disks: reduce seek time

16.7 Summary

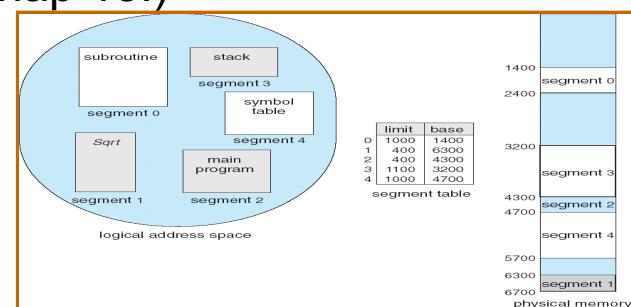
Segmentation

- ✓ Divide address space into logical regions called **segment**
- ✓ Overcome the memory wasted between segments (e.g. heap and stack in the base/limit mechanism)
- ✓ Flexible: **support sharing and protection**

But, still have some problems

- ✗ Variable size **relatively hard to implement in hardware**, may cause **external fragmentation** which complicate free space management
- ✓ Memory waste within a segment, especially **sparse segment** **need to allocate address space that are actually used by a process**

✗ Alternative: **fixed size** **Paging** (chap 18.)



Chap. 17 Free Space Management

Free-space management

→ Variable size (e.g. malloc() or segmentation)

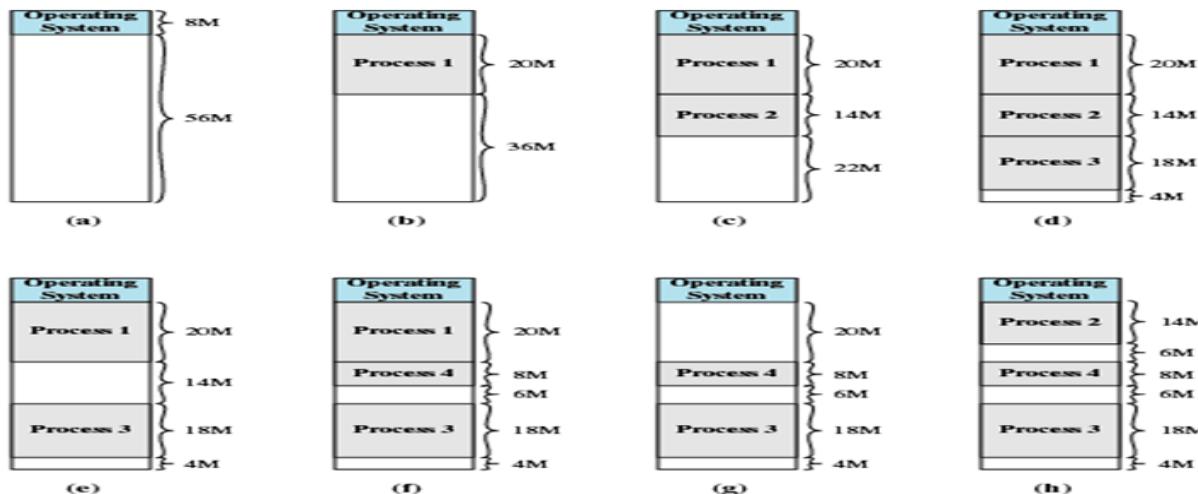
Complicate, need to handle external fragmentation [?]

✓ Fixed size (e.g. paging)

Relatively easy, usually a list of free fixed-size units [?] later chapters

Internel

dbz



(Source: A. Silberschatz, "Operating system Concept")

Process 2 is "relocated" dynamically

Need the swap space (in a disk) when a process is suspended.

How to handle when a new process is forked at (h) step whose size is 3 or 10MB?

17.1 Assumptions

~~Interfaces~~

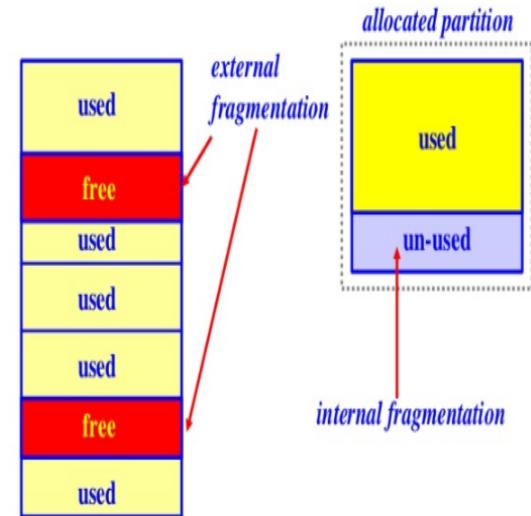
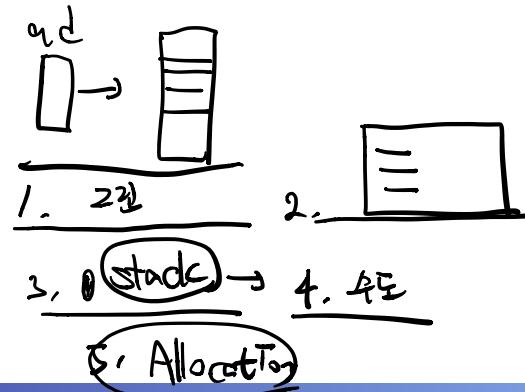
✓ ~~malloc()/free()~~

Free space

- ✓ Managed by a list (**free list**)
- ✓ In actual OSes, free space is managed by various data structures including a hashed list or tree (e.g. buddy system)

Fragmentation

- ✓ External: variable-size allocation
- ✓ Internal: fixed-size allocation
- ✓ Focus on external fragmentation



17.2 Low-level Mechanisms

Splitting and Coalescing

- ✓ Memory: 30-byte heap

- ✓ Free list

- ✓ Request

10B \square alloc

Larger than

Smaller than 10B \square need splitting

- Allocate 1 byte

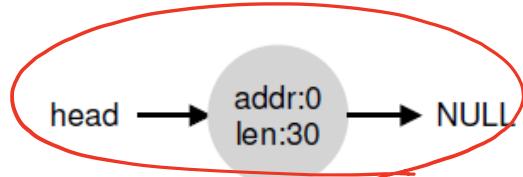
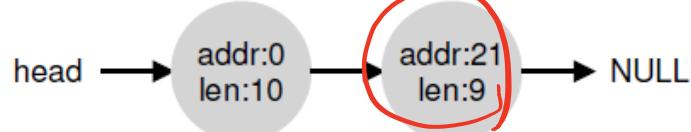
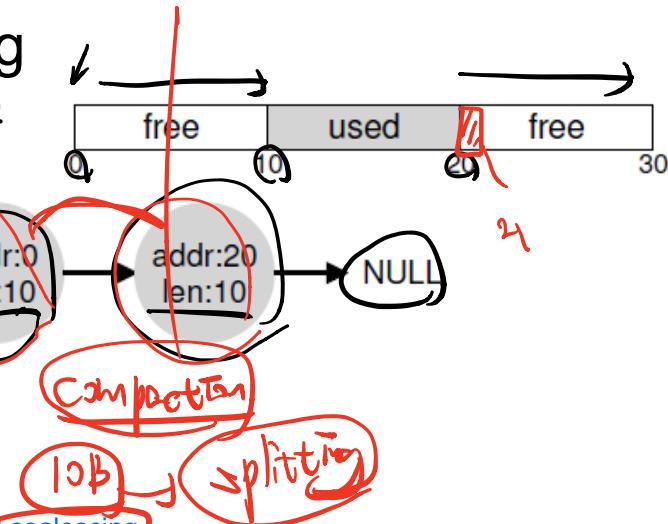
- ✓ Free

Free the used space 10~19 \square need coalescing

- Sort free entries, check neighbors when inserting into the free list

이전에
다 알아두기

이전
페이지



17.3 Basic Strategies

Free-space allocation policy

Best-fit

allocate from the smallest chuck which is bigger than the request size

Worst-fit

allocate from the largest chuck which is bigger than the request size

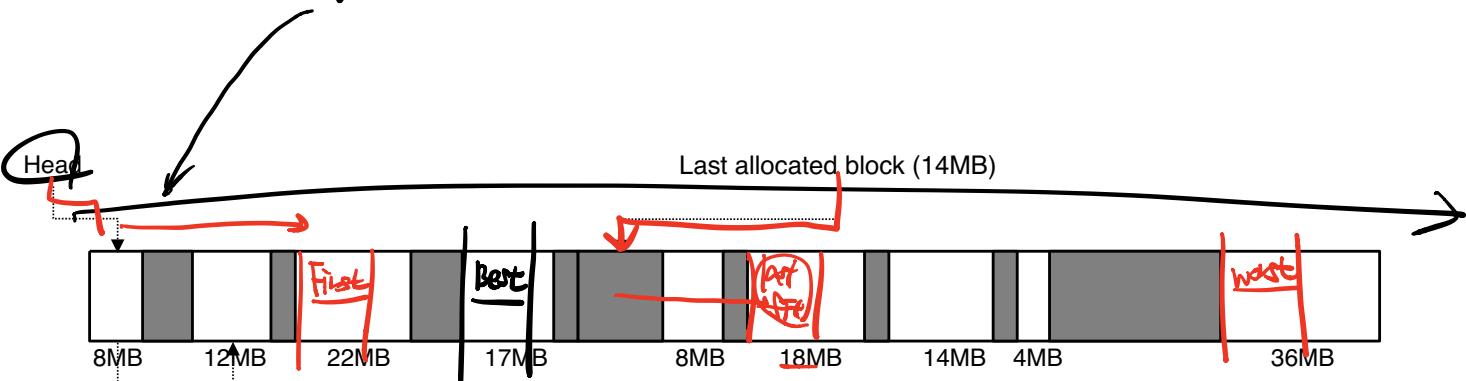
First-fit

allocate from the first chuck which is bigger than the request size, search start from head

Next-fit

allocate from the first chuck which is bigger than the request size, search start from the last allocated chunk

21 25 15 31 25 10 5



Need to allocate 16MB available space. Which one by each policy?

17.4 Other Approaches

Buddy allocation

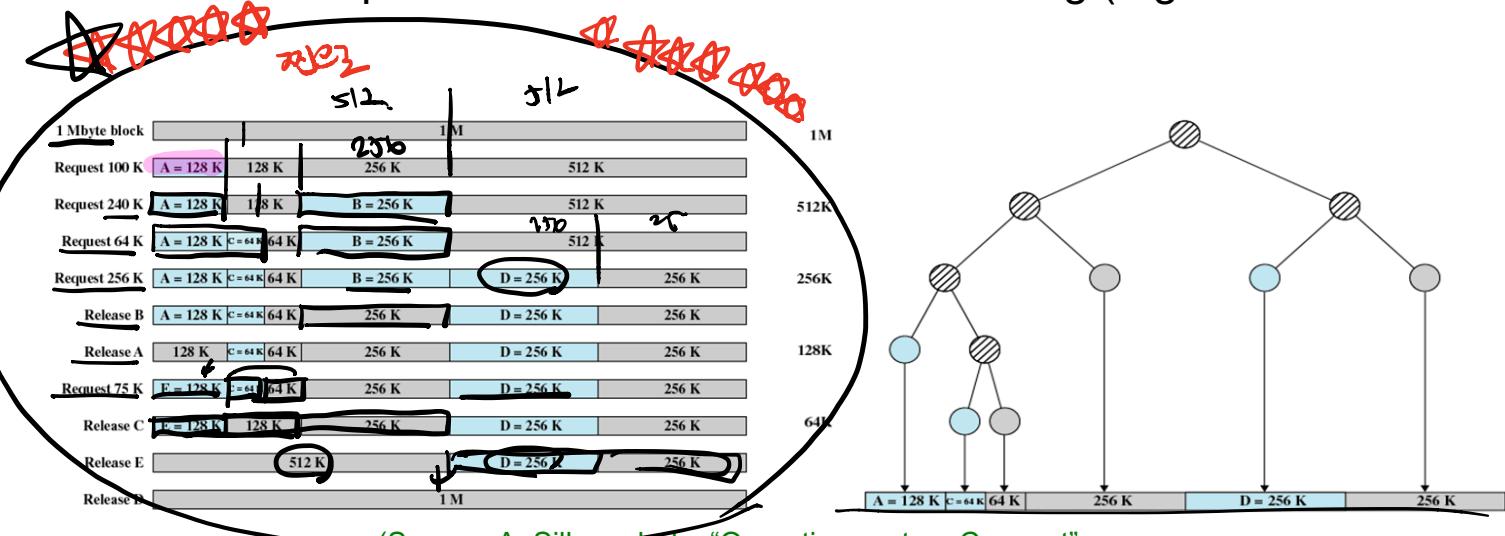
- ✓ To make splitting/coalescing simple
- ✓ Allocate a free memory with the size of 2^n (e.g. 4KB, 8KB, ...)

Segregated Lists

- ✓ Some applications have one (or a few) popular-sized request
- ✓ Manage them in a segregated list \square same size \square easier to split and coalescing
- ✓ Popular example: **slab** allocator in Solaris (and in Linux)

Others

- ✓ More complex data structure for fast searching (e.g. balanced B-tree)



(Source: A. Silberschatz, "Operating system Concept")

17.5 Summary

Memory virtualization: Goal

- ✓ Transparency, isolation, efficiency
- ✓ Address space: Isolated, private-owned memory
- ✓ Address translation: virtual to physical address

Dynamic relocation

- ✓ Base & Limit (Bound) approach
- ✓ Generalized approach \square segmentation

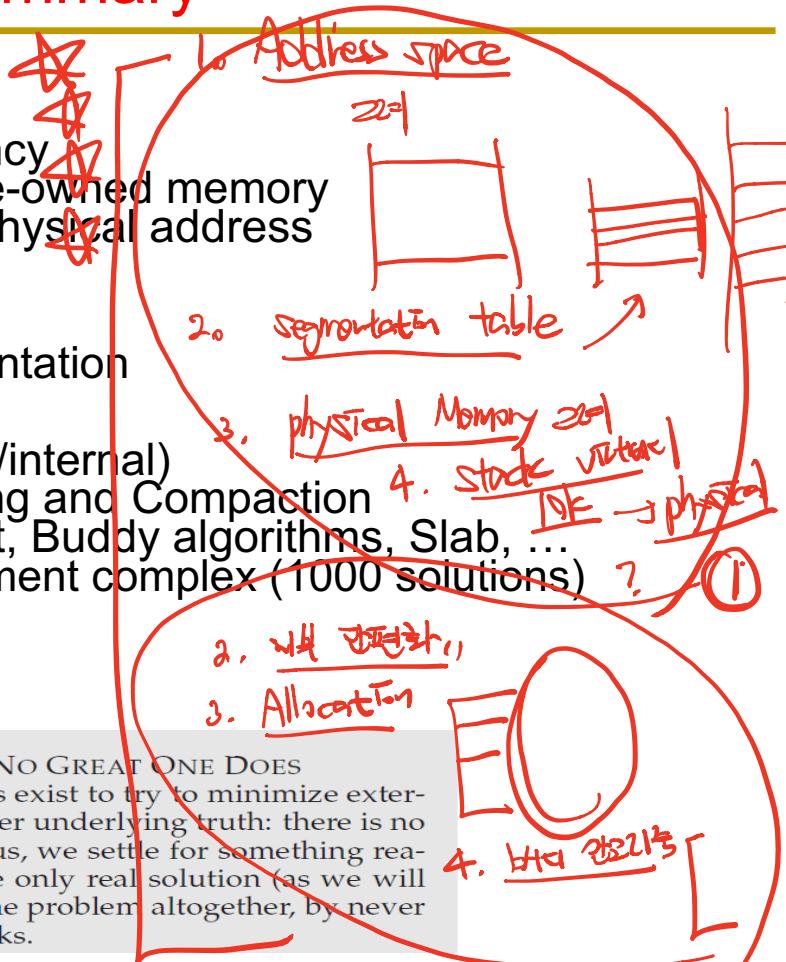
Free-Space Management

- ✓ Reduce fragmentation (external/internal)
- ✓ Mechanism: Splitting, Coalescing and Compaction
- ✓ Policy: Best fit, First fit, Worst fit, Buddy algorithms, Slab, ...
- ✓ \square Variable size makes management complex (1000 solutions) ?

TIP: IF 1000 SOLUTIONS EXIST, NO GREAT ONE DOES

The fact that so many different algorithms exist to try to minimize external fragmentation is indicative of a stronger underlying truth: there is no one "best" way to solve the problem. Thus, we settle for something reasonable and hope it is good enough. The only real solution (as we will see in forthcoming chapters) is to avoid the problem altogether, by never allocating memory in variable-sized chunks.

- Simple question to take attendance: 1) Among the topics you have learnt in LN 8, what is the most interesting topic for you? Explain your opinion with some reasonable grounds. 2) Suggest an exam question about the topic you choose (it may appear in the final exam \square) (until 6 PM, June 10th).



Appendix: 17.2 Low-level Mechanisms

Tracking the size of allocated regions

- ✓ `free()`: argument **?** pointer only, not size

Need to track the size of a unit that is freed for coalescing

Most allocators utilize a **header** block, usually just before the handed-out chunk of memory

- Size, magic number for integrity checking, additional pointer to speed up deallocation, and other information

```
typedef struct __header_t {  
    int size;  
    int magic;  
} header_t;
```

```
void free(void *ptr) {  
    header_t *hptr = (void *)ptr - sizeof(header_t);  
    ...
```

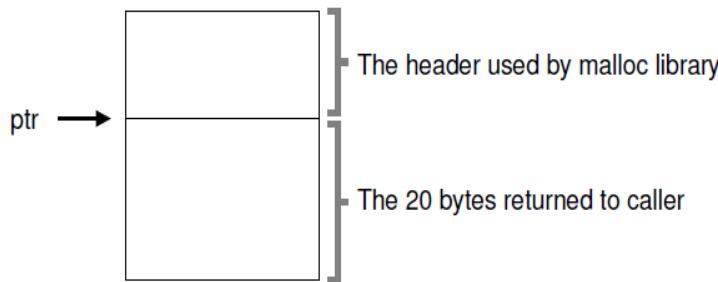


Figure 17.1: An Allocated Region Plus Header

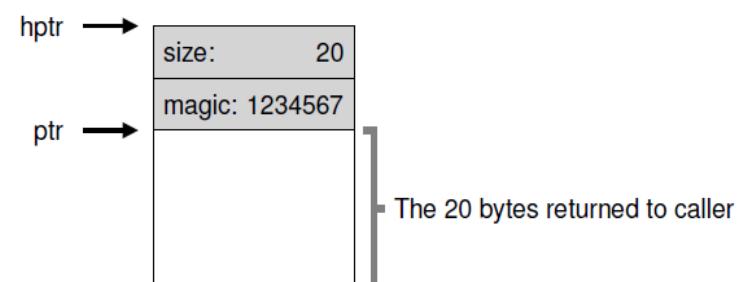


Figure 17.2: Specific Contents Of The Header

Appendix: 17.2 Low-level Mechanisms

Embedding the free list into a heap

- ✓ Figure 17.3: initial stage, build a free list inside the free space
Free space: 4KB (4096 byte), entry of the free list: 8 byte (size, next) \square size becomes 4088.
- ✓ Figure 17.4: after “malloc(100)”
Header for the allocated space: 8 byte (size, magic #) \square 3980 (split occurs)
Head: pointer for the free list, ptr: pointer returned to malloc()
- ✓ Figure 17.5: after three “malloc(100)”s \square 3764

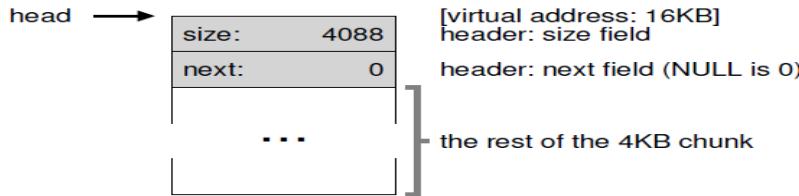


Figure 17.3: A Heap With One Free Chunk

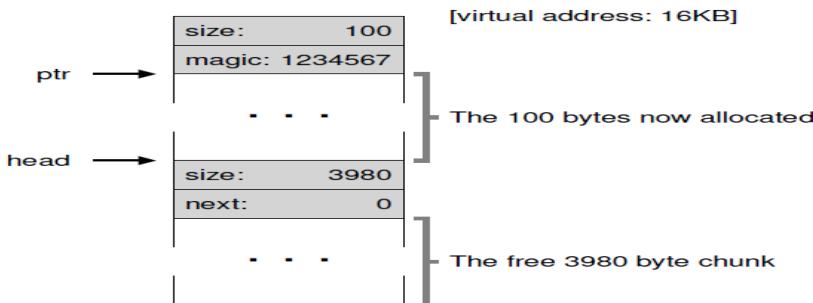


Figure 17.4: A Heap: After One Allocation

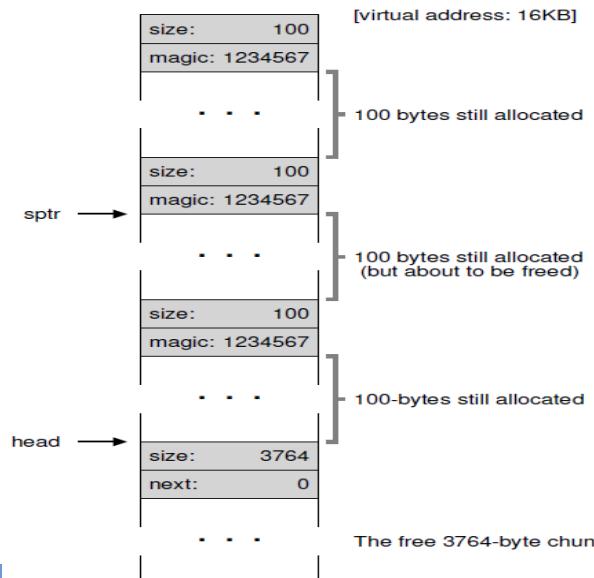


Figure 17.5: Free Space With Three Chunks Allocated

Appendix: 17.2 Low-level Mechanisms

Embedding the free list into a heap

- ✓ Figure 17.5: after three “malloc()”s, trigger one “free(sptr)” request
- ✓ Figure 17.6: after “free(sptr)”
 - Two entries in the free list: head ↗ (100, 16708) ↗ (3764, 0 (NULL))
 - Virtual address 16708 = $16 \times 1024 + 3 \times 108$
- ✓ Figure 17.7: after three “free()”s
 - Compaction-less version (c.f. Compaction version: Figure 17.3)

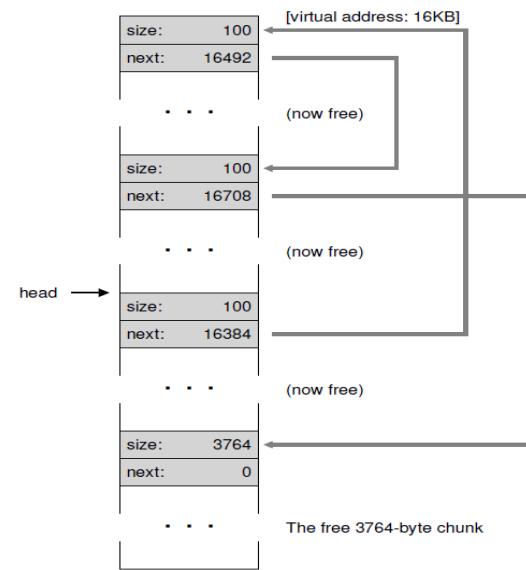
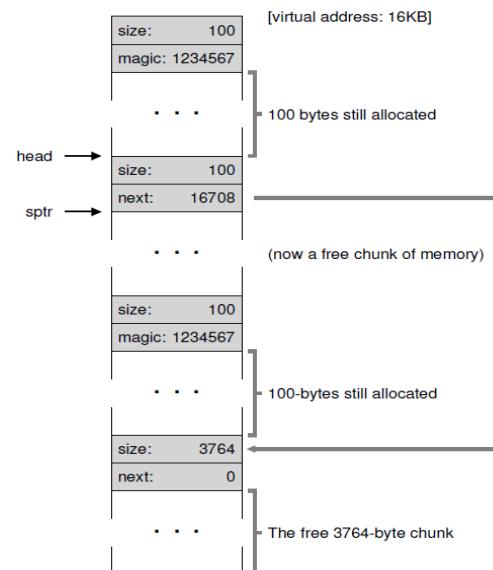
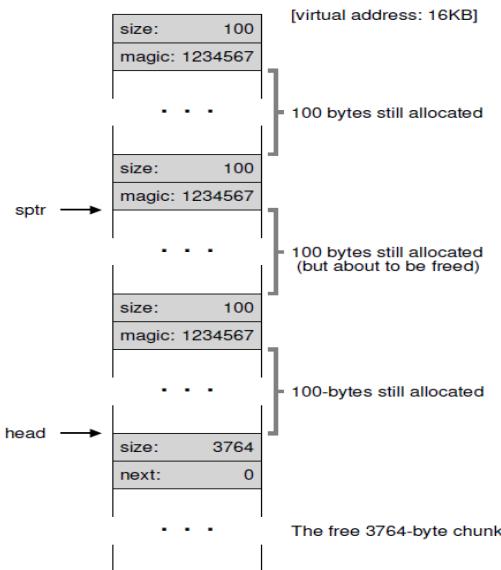


Figure 17.5: Free Space With Three Chunks Allocated

Figure 17.6: Free Space With Two Chunks Allocated

Figure 17.7: A Non-Coalesced Free List